Ensemble Learning

Training individual learners & combining their predictions

CS 6316 – Machine Learning Fall 2017

OUTLINE

- Motivation for combining methods
- Ensemble Learning
- Example: Random Forests

Motivation for Combining Methods

- General setting (used in this course)
 - Given training data set
 - Flexible model parameterization ~ learning method
 - Empirical loss function
 - Optimization method
- Select the best model via single application of a learning method to data

Learning Method + Data → Predictive Model

Motivation (cont'd)

Learning Method + Data → Predictive Mode

- Theoretical and empirical evidence
 - No single "best" method exists
- Always possible to find:
 - Best method for given data set
 - Best data set for given method
- Many philosophical approaches (eastern philosophy / Bayesian averaging/etc) combine several theories (models) explaining the data

Collective Decision Making

Commonly used in our daily lives

- Jury trial
- Multiple expert opinions (in medicine, law, ...)



Collective Decision Making

When it works

- Experts are indeed intelligent
- Expert' opinions are different (not correlated)
- Their decisions are combined intelligently

• When it does not work

- Majority are not "smart"
- Experts give similar opinions (highly correlated)
- Their decisions are not combined intelligently

Strategies for Combining Methods

- Standard inductive learning setting
- Two combining strategies (for improved generalization)
 - 1. Apply different learning methods to the same data
 - → Committee of Networks, Stacking, Bayesian averaging
 - 2. Apply the same method to different (modified) realizations of training data
 - → Bagging, Boosting

Strategies for Combining Methods

Combining methods are used for

Classification:

$$F(\mathbf{x}) = sign\left(\sum_{k=1}^{N} w_k f_k(\mathbf{x})\right)$$

And **Regression**:

$$F(\mathbf{x}) = \sum_{k=1}^{N} w_k f_k(\mathbf{x})$$

What and Why: Ensemble Learning

What is ensemble learning?

 Ensemble learning refers to a collection of methods that learn a target function by training a number of individual learners and combining their predictions

• Why ensemble learning?

- Accuracy: a more reliable mapping can be obtained by combining the output of multiple "experts"
- Efficiency: [divide-and conquer approach] a complex problem can be decomposed into multiple subproblems that are easier to understand and solve

Ensemble Learning

- There is not a single model that works for all problems!
- "To solve really hard problems, we'll have to use several different representations..... It is time to stop arguing over which type of technique is best..... Instead we should work at a higher level of organization and discover how to build managerial systems to exploit the different virtues and evade the different limitations of each of these ways of comparing things." [Minsky, 1991]

Ensemble Learning

• When to use ensemble learning?

- When you can build component classifiers that are more accurate than chance and,
- more importantly, that are independent from each other

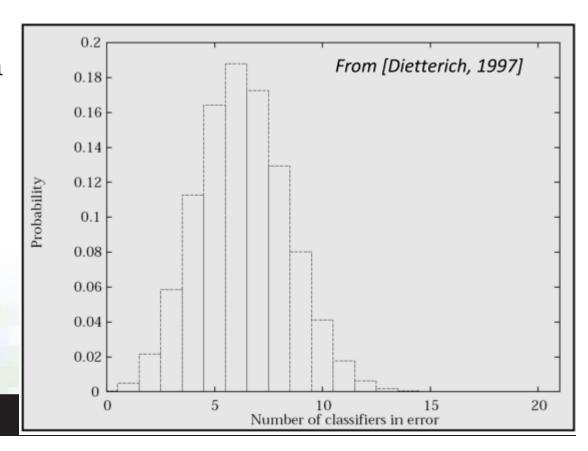
Why do ensembles work?

Because uncorrelated errors of individual classifiers can be eliminated through averaging

- Assume a binary classification problem for which you can train individual classifiers with **an error rate of 0.3**
- Assume that you **build an ensemble** by combining the prediction of 21 such classifiers **with a majority vote**
- What is the probability of error for the ensemble?

Why do ensembles work?

- What is the probability of error for the ensemble?
 - In order for the ensemble to misclassify an example, 11 or more classifiers have to be in error, or a probability of 0.026
 - The histogram here shows the distribution of the number of classifiers that are in error in the ensemble machine

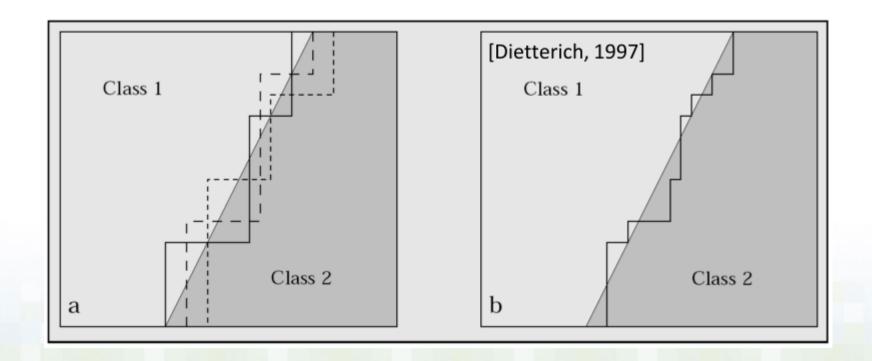


Ensemble Learning

- The target function may not be implementable with single classifiers, but may be approximated by ensemble averaging
 - Assume that you want to build a diagonal decision boundary with decision trees
 - The decision boundaries constructed by these machines are hyperplanes parallel to the coordinate axes, or "staircases" in the example (next slide)

Ensemble Learning

• By averaging a large number of such "staircases", the diagonal decision boundary can be approximated with arbitrarily small accuracy



Tree Based Ensemble Methods

- Ensemble methods involve group of predictive models to achieve a better accuracy and model stability
- Ensemble methods are known to impart supreme boost to tree based models
- Like every other model, a tree based model also suffers from the plague of bias and variance
 - Bias: how much on an average are the predicted values different from the actual value
 - Variance: how different will the predictions of the model be at the same point if different samples are taken from the same population

Tree Based Ensemble Methods

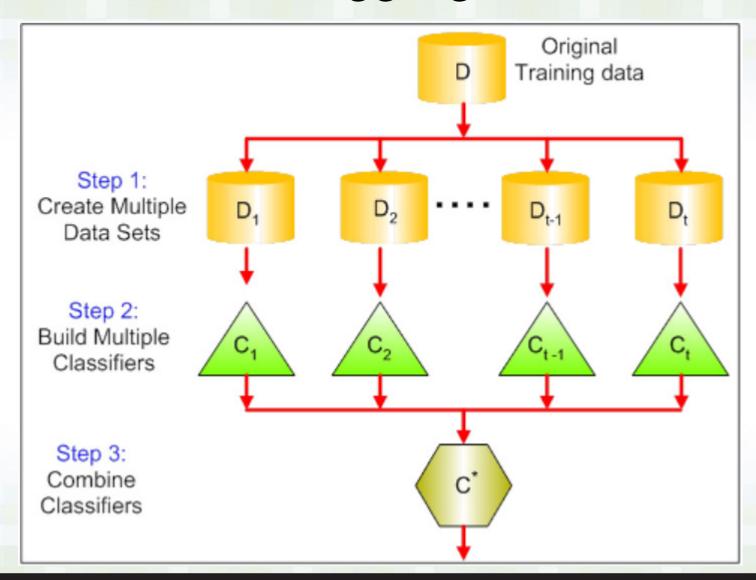
- If you have a small tree → you will get a model with low variance and high bias
- As you increase the complexity of your model, you will see a reduction in prediction error due to lower bias in the model
- As you continue to make your model more complex, you end up over-fitting your model and your model will start suffering from **high variance**

Tree Based Ensemble Methods

- An optimal model should maintain a balance between these two types of errors. This is known as the trade-off management of bias-variance errors
- Ensemble learning is one way to execute this trade off analysis

What is it? How does it work?

• Bagging is a technique used to reduce the variance of our predictions by combining the result of multiple classifiers modeled on different sub-samples of the same data set



• Create Multiple Data sets:

- Sampling is done with replacement on the original data and new datasets are formed
- The new data sets can have a fraction of the columns as well as rows, which are generally hyper-parameters in a bagging model
- Taking row and column fractions helps in making robust models, less prone to overfitting

• Build Multiple Classifiers:

- Classifiers are built on each data set
- Generally the same classifier is modeled on each data set and predictions are made

Combine Classifiers:

- The predictions of all the classifiers are combined using a mean, median or mode value depending on the problem at hand
- The combined values are generally more robust than a single model

• The number of models built is not a hyper-parameter*. Higher number of models are always better or may give similar performance than lower numbers. It can be theoretically shown that the variance of the combined predictions are reduced to 1/n (n: number of classifiers) of the original variance, under some assumptions

*hyperparameters are parameters whose values are set prior to the commencement of the learning process

Bagging // Random Forest

- There are various implementations of bagging models
- Random forest is one of them



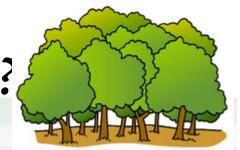


Random Forest



- Random Forest is a **versatile** machine learning method capable of performing both regression and classification tasks. It also undertakes dimensional reduction methods, handles outlier values and other essential steps of data exploration, and does a fairly good job
- It is a type of ensemble learning method, where a group of weak models combine to form a powerful model

RF: How does it work?



- In Random Forest (RF), multiple trees are grown (hence "forest") as opposed to a single tree
- To classify a new object based on attributes, each tree gives a classification (the tree "votes" for that class)
- The forest chooses the classification having the most votes (over all the trees in the forest)
- In the case of **regression**, it takes the average of outputs by different trees

Random Forest



- Assume number of cases in the training set is N. Then, sample of these N cases is taken at random but with replacement. This sample will be the training set for growing the tree
- If there are M input variables, a number *m*<M is specified such that at each node, *m* variables are selected at random out of the M. The best split on these *m* is used to split the node. The value of *m* is held constant while we grow the forest
- Each tree is grown to the largest extent possible and there is no pruning
- Predict new data by **aggregating** the predictions of the ntree trees (i.e., majority votes for **classification**, average for **regression**)

Voting



- So what good are 10000 (*probably*) **bad models**? Well it turns out that they really aren't that helpful. But what is helpful are the few really good decision trees that you also generated along with the bad ones
- When you make a prediction, the new observation gets pushed down each decision tree and assigned a predicted value/label. Once each of the trees in the forest have reported its predicted value/label, the predictions are tallied up and the mode vote of all trees is returned as the final prediction

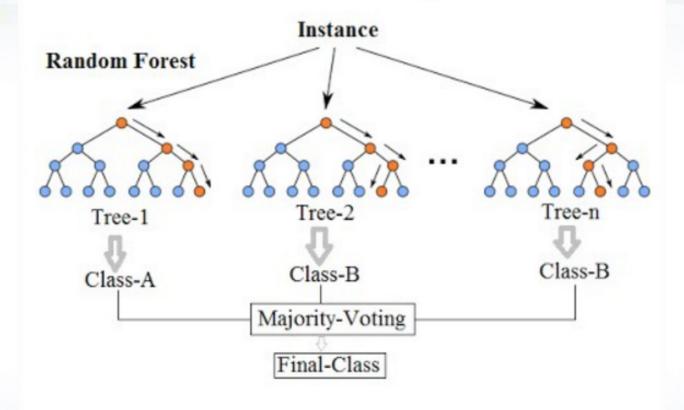
Voting

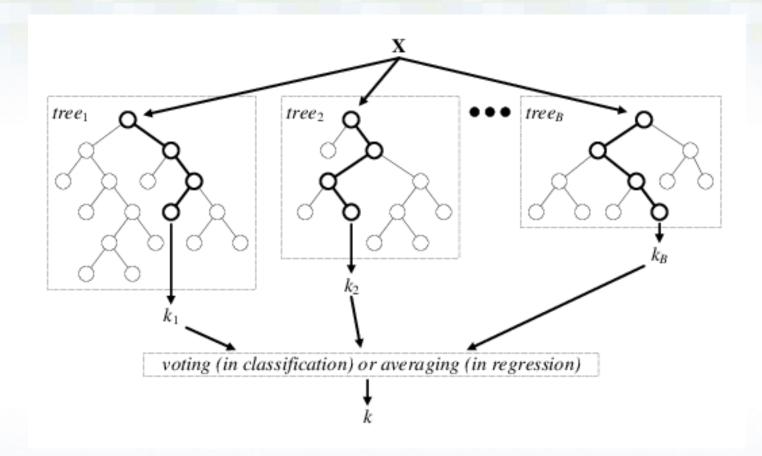


• Simply, the 99.9% of trees that are *irrelevant* make predictions that are all over the map and cancel each another out. The predictions of the minority of trees that are good top that noise, and yield a good prediction



Random Forest Simplified





Advantages of RF

- Suitable for both classification and regression
- Can handle large data sets with a large number of attributes (dimensions) ~ 1000s!
- Can assist with dimensionality reduction because often implementations help to identify the most significant variables
 - Can output importance of variables, which can be handy when exploring new/unknown data sets

Advantages of RF

- It has an effective method for estimating missing data and maintains accuracy when a large proportion of the data are missing
- It has methods for balancing errors in data sets where classes are imbalanced

Note

- RF involves sample of the input data with replacement
 - This is called *bootstrap sampling*
- Here one third (for example) of the data is *not* used for training and can be used for testing. These are called the **out of bag** samples. Error estimated on these out of bag samples is known as *out of bag error*. Studies have shown that error estimates by Out of bag is as accurate as using a test set of the same size as the training set.
- Therefore, using the out-of-bag error estimate removes the need for a set aside test set

Out-of-Bag

• When the training set for the current tree is drawn by sampling with replacement, about one-third of the cases are left out of the sample. This *oob* (out-of-bag) data is used to get a running unbiased estimate of the classification error as trees are added to the forest. It is also used to get estimates of variable importance

Out-of-Bag

- So, for RF, there is *no need* for cross-validation or a separate test set to get an unbiased estimate of the test set error. It is estimated internally
 - Each tree is constructed using a different bootstrap sample from the original data
 - About one-third of the cases are left out of the bootstrap sample and not used in the construction of the kth tree

Disadvantages of RF

- While it does a good job with classification problems, it doesn't perform as well with regression problems
- In the case of regression, it doesn't predict beyond the range in the training data
 - May over fit data sets that are particularly noisy

References:

- ~Ensemble Learning ~ R. Gutierrez-Osuna ~ Pattern Analysis ~ TAMU
- ~Tree Based Modeling from Scratch ~ Analytics Vidhya
- ~Random Forests ~ L.Breiman and A.Cutler ~ stat.berkeley.edu