

Ensemble Learning

Ensembles, Bagging, Boosting, AdaBoost Technique

CS 6316 – Machine Learning

Fall 2017

OUTLINE

- Methods for Constructing Ensembles
- Combination Strategies
- Mixtures of Experts (ME)
- Bagging
- Boosting
- AdaBoost
- AdaBoost Steps
- AdaBoost Example

Methods for Constructing Ensembles


- **Subsampling the training examples**
 - Multiple hypotheses are generated by training individual classifiers on different datasets obtained by resampling a common training set (Bagging, Boosting)
- **Manipulating the input features**
 - Multiple hypotheses are generated by training individual classifiers on different representations, or different subsets of a common feature vector

Methods for Constructing Ensembles

- **Manipulating the output targets**
 - The output targets for C classes are encoded with an L -bit codeword, and an individual classifier is built to predict each one of the bits in the codeword
 - Additional “auxiliary” targets may be used to differentiate classifiers
- **Modifying the learning parameters of the classifier**
 - A number of classifiers are built with different learning parameters, such as number of neighbors in a k Nearest Neighbor rule, initial weights in an MLP, etc.

Structure of Ensemble Classifiers (Parallel)

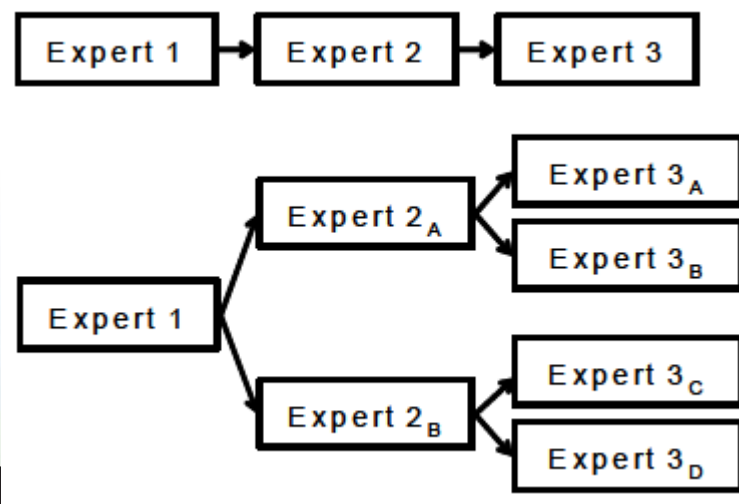
- All the **individual classifiers are invoked independently**, and **their results are fused with a combination rule** (e.g., average, weighted voting) or a meta-classifier (e.g., stacked generalization)
- The majority of ensemble approaches in the literature fall under this category



Jain, 2000

Structure of Ensemble Classifiers (Cascading or Hierarchical)

- Classifiers are **invoked in a sequential** or **tree-structured** fashion
- For the purpose of efficiency, **inaccurate but fast methods are invoked first** (maybe using a small subset of the features), and **computationally more intensive but accurate methods are left for the latter stages**



Jain, 2000

Combination Strategies (Static Combiners)


- The combiner decision rule is independent of the feature vector. Static approaches can be broadly divided into non-trainable and trainable:
- **Non-trainable:** The voting is performed independently of the performance of each individual classifier
- **Trainable:** The combiner undergoes a separate training phase to improve the performance of the ensemble

Non-trainable

- Various combiners may be used, depending on the type of output produced by the classifier, including:
 - **VOTING**: used when each classifier produces a single class label. In this case, each classifier “votes” for a particular class, and the class with the majority vote on the ensemble wins
 - **AVERAGING**: used when each classifier produces a confidence estimate (e.g., a posterior). In this case, the winner is the class with the highest average posterior across the ensemble
 - **BORDA COUNTS**: used when each classifier produces a rank. The Borda count of a class is the number of classes ranked below it [Ho et al., 1994]

Trainable

- **Weighted averaging:** the output of each classifier is weighted by a measure of its own performance, e.g., prediction accuracy on a separate validation set
- **Stacked generalization:** the output of the ensemble serves as a feature vector (input) to a meta-classifier (second-level expert)

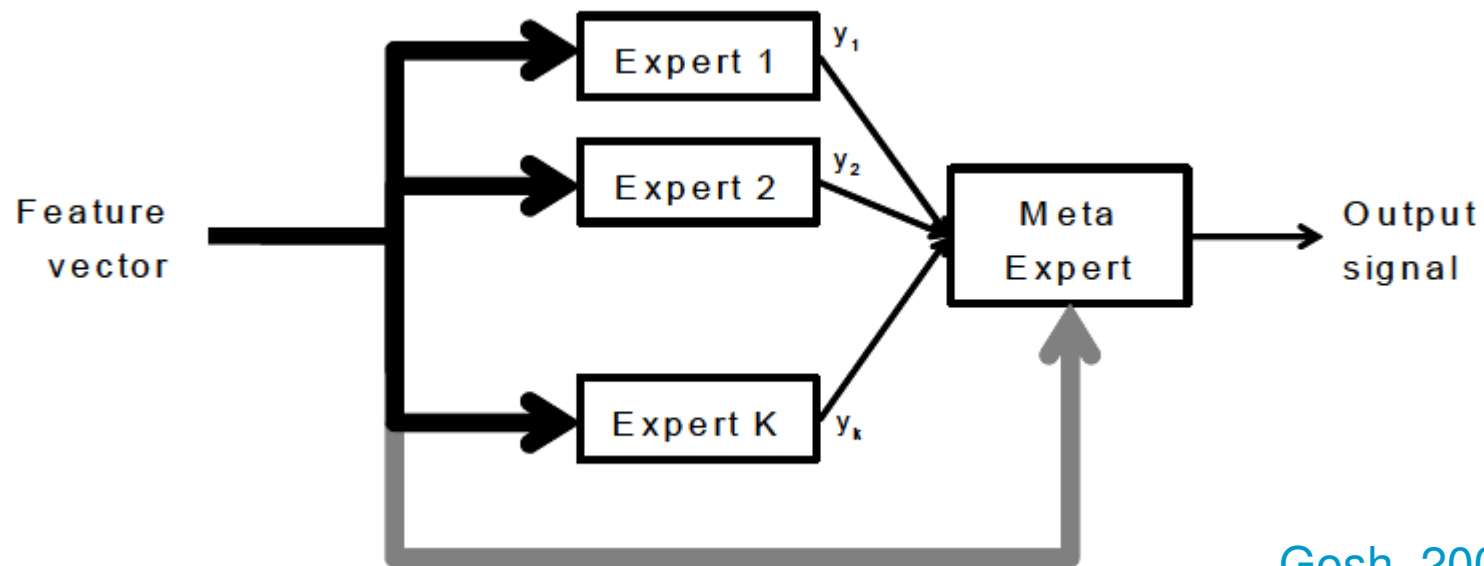


Adaptive Combiners

- The combiner is a function that depends on the input feature vector
 - Thus, the ensemble implements a function that is local to each region in feature space
 - This divide-and-conquer approach leads to modular ensembles where relatively simple classifiers specialize in different parts of I/O space
 - In contrast with static-combiner ensembles, the individual experts here do not need to perform well for all inputs, only in their region of expertise

Adaptive Combiners

- Representative examples of this approach are Mixture of Experts (ME) and Hierarchical ME [Jacobs et al., 1991; Jordan and Jacobs, 1994]

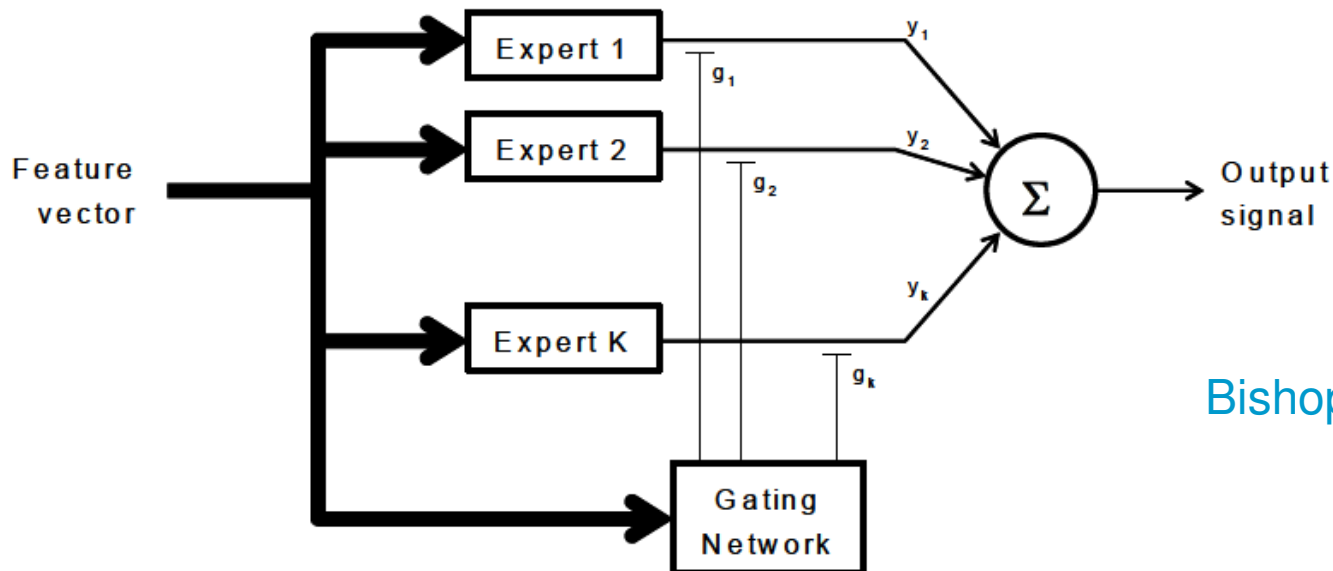


Gosh, 2002

Mixture of Experts (ME)

ME is the classical adaptive ensemble method

- A gating network is used to partition feature space into different regions, with one expert in the ensemble being responsible for generating the correct output within that region [Jacobs et al., 1991]
- The experts in the ensemble and the gating network are trained simultaneously



Bishop, 1995

Bagging

- Subsampling the training set → Bagging [Breiman, 1996]
- Bagging (for **b**ootstrap **a**ggregation) creates an ensemble by **training individual classifiers on bootstrap samples** of the training set
- As a result of the **sampling-with-replacement** procedure, **each classifier is trained on the average of 63.2% of the training examples**
- Bagging traditionally uses **component classifiers of the same type** (e.g., decision trees), and a simple combiner consisting of a majority vote across the ensemble

Boosting

- [Schapire, 1990; Freund and Schapire, 1996]
- Boosting takes a different resampling approach than bagging, which maintains a constant probability of $1/N$ for selecting each example
- In boosting, this probability is adapted over time based on performance
 - The component classifiers are built sequentially, and examples that are mislabeled by previous components are chosen more often than those that are correctly classified

Boosting

- Boosting is based on the concept of a “*weak learner*”, an algorithm that performs slightly better than chance (e.g., 50% classification rate on binary tasks)
 - Schapire has shown that a weak learner can be converted into a strong learner by changing the distribution of training examples
 - *Small* benefits achieved by using *highly accurate* classifiers
 - There are a number of variants of boosting available in literature
 - A popular one is **AdaBoost** (Adaptive Boosting) – which allows the designer to continue adding components until an arbitrarily small error rate is obtained on the training set

AdaBoost

- AdaBoost (Adaptive Boosting) is a popular boosting technique which helps combine multiple “weak classifiers” into a single “strong classifier”
- A weak classifier is simply a classifier that performs poorly, but **performs better than random guessing**
- *A simple example might be classifying a person as male or female based on their **height**. You could say **anyone over 5' 9"** is a **male** and anyone under that is a **female**. You'll misclassify a lot of people that way, **but your accuracy will still be greater than 50%***

AdaBoost

- AdaBoost can be applied to any classification algorithm, so it's really a technique that **builds on top of other classifiers** as opposed to being a classifier itself

AdaBoost: What Do We Get?

- You could just train a bunch of weak classifiers on your own and combine the results, **so what does AdaBoost do for you?** There's really two things it figures out for you:
 1. It helps you **choose the training set for each new classifier** that you train based on the results of the previous classifier
 2. It **determines how much weight should be given** to each classifier's proposed answer when combining the results

AdaBoost: Training Set Selection

- Each weak classifier should be trained on a random subset of the total training set
- The subsets can overlap—it's not the same as, for example, dividing the training set into ten portions
- AdaBoost assigns a “weight” to each training example, which determines the probability that each example should appear in the training set
- Examples with higher weights are more likely to be included in the training set, and vice versa

AdaBoost: Training Set Selection

- *Examples with higher weights are more likely to be included in the training set, and vice versa*
- After training a classifier, **AdaBoost increases the weight on the misclassified examples** so that these examples will make up a larger part of the next classifiers training set, and **hopefully the next classifier trained will perform better on them**

AdaBoost

- Assume 2-class problem, with labels +1 and -1
– y^i in $\{-1, 1\}$

- Discriminant function:

$$g(x) = \sum_{t=1}^T \alpha_t h_t(x) = \alpha_1 h_1(x) + \alpha_2 h_2(x) + \dots + \alpha_T h_T(x)$$

- Where $h_t(x)$ is a weak classifier, for example:

$$h_t(x) = \begin{cases} -1 & \text{e.g. if email has word "money" } \Rightarrow \text{Spam} \\ 1 & \text{e.g. if email doesn't have word "money" } \Rightarrow \text{Not Spam} \end{cases}$$

- The final classifier is the sign of the discriminant function

$$f_{final}(x) = \text{sign} \left(\sum_t \alpha_t h_t(x) \right)$$

Idea Behind AdaBoost

- Algorithm is **iterative**
- Maintains distribution of weights over the training examples
- Initially weights are **equal**
- Main Idea: at successive iterations, the **weight of misclassified examples is increased**
- This **forces the algorithm to concentrate on examples that have not been classified correctly so far**

Idea Behind AdaBoost

- Examples of high weight are shown more often at later rounds
- Face/nonface classification problem:

Round 1

best weak classifier:

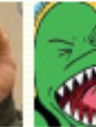
change weights:

						
$1/7$	$1/7$	$1/7$	$1/7$	$1/7$	$1/7$	$1/7$
✓	✗	✓	✓	✗	✓	✗
$1/16$	$1/4$	$1/16$	$1/16$	$1/4$	$1/16$	$1/4$

Round 2


best weak classifier:


change weights:

									
✓	✓	✓	✗	✗	✗	✓	✓	✓	✓
$1/8$	$1/32$	$1/32$	$11/32$		$1/2$		$1/8$	$1/32$	$1/32$

Idea Behind AdaBoost

Round 3



- out of all available weak classifiers, we choose the one that works best on the data we have at round 3
- we assume there is always a weak classifier better than random (better than 50% error)
-  image is half of the data given to the classifier
- chosen weak classifier **has to** classify this image correctly

Additional Comments

- Ada boost is very simple to implement, provided you have an implementation of a “weak learner”
- Will work as long as the “basic” classifier $h_t(x)$ is at least **slightly better than random**
 - will work if the error rate of $h_t(x)$ is **less than 0.5**
 - 0.5 is the error rate of a *random guessing* for a 2-class problem

AdaBoost for 2 Classes

- **Initialization step:** for each example x , set weight

$$D_1(x) = \frac{1}{N}, \text{ where } N \text{ is the number of examples}$$

- **Iteration step** (for $t = 1 \dots T$):
 1. Find best weak classifier $h_t(x)$ using weights $D(x)$ of ex's
 2. Compute the error rate ϵ_t as

$$\epsilon_t = \sum_{i=1}^N D(x^i) \cdot \underbrace{I[y^i \neq h_t(x^i)]}$$

*Will only accumulate weights of examples that were **misclassified** (i.e. $y^i \neq h_t(x^i)$)*

$$= \begin{cases} 1 & \text{if } y^i \neq h_t(x^i) \\ 0 & \text{otherwise} \end{cases}$$

AdaBoost

- Iteration step (for $t = 1 \dots T$) (CONTINUED)

3. Compute weight α_t of classifier h_t

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_t}{\varepsilon_t} \right) > 0 \quad \text{Note: natural log}$$

4. For each example x^i , update the weights ($t+1$ iteration):

$$D_{t+1}(x^i) = D_t(x^i) \cdot \exp(\alpha_t \cdot I[y^i \neq h_t(x^i)])$$

Recall: $\exp(x) = e^x$

Normalize $D_{t+1}(x^i)$ so that $\sum_{i=1}^N D(x^i) = 1$

$$D_{t+1}(x^i) = \frac{D_t(x^i) \cdot \exp(\alpha_t \cdot I[y^i \neq h_t(x^i)])}{Z_t} \quad \leftarrow \text{Normalization constant}$$

AdaBoost

- The final combined classifier that will classify example “x”

$$H_{final}(x) = \text{sign} \left(\sum_t \alpha_t h_t(x) \right)$$


AdaBoost

$$H_{final}(x) = \text{sign} \left(\sum_t \alpha_t h_t(x) \right)$$

- The final classifier consists of “T” weak classifiers
- $h_t(x)$ is the output of weak classifier \mathbf{t} (-1/+1 in this ex.)
- α_t is the weight applied to classifier \mathbf{t} as determined by AdaBoost
- So the final output is just a linear combination of all of the weak classifiers. **Final decision: LOOKING AT SIGN OF THE SUM**


A little about... α_t

- Plot of what α_t will look like for classifiers with different error rates:



A little about... α_t


1. The classifier weight grows exponentially as the error approaches 0 (better classifiers are given exponentially more weight)
2. The classifier weight is zero if the error rate is 0.5.



A classifier with 50% accuracy is no better than random guessing, so we ignore it!

3. The classifier weight grows exponentially negative as the error approaches 1.

A little about... α_t




3. The classifier weight grows exponentially negative as the error approaches 1. We give negative weight to classifiers with worse than 50% accuracy
→ *“Whatever the classifier says, do the opposite!”*

Recall: $\exp(x) = e^x$ $e = \text{Euler constant}$

Exp() term when updating weights

- In the equation to **update the weights (D)** there exists the **exp** term. $\exp(x)$ behaves as follows:

*Exp(x) will return a fraction for **NEGATIVE** values of x*



*Exp(x) will return a value greater than one for **POSITIVE** values of x*

So the weight for training sample i will be either increased or decreased depending on the final sign of the term


AdaBoost: Step 1 [**important**]

[1] Find best weak classifier $h_t(x)$ using weights $D(x)$

- Some classifiers **accept weighted samples**, but many don't
- If classifier does not take weighted samples, **sample from the training samples according to the distribution $D(x)$**




- Draw k samples, each x with probability equal to $D(x)$:



AdaBoost: Step 1

[1] Find best weak classifier $h_t(x)$ using weights $D(x)$

- Give to the classifier the re-sampled examples:



- To find the best weak classifier, go through **all** weak classifiers, and find the one that gives the **smallest error** on the re-sampled examples








Weak Classifiers	$h_1(x)$	$h_2(x)$	$h_3(x)$...	$h_m(x)$
Errors:	0.46	0.36	0.16	...	0.43

The best classifier to choose at iteration t

AdaBoost: Step 2

[2] Compute ε_t the error rate

$$\varepsilon_t = \sum_{i=1}^N D(x^i) \cdot \underbrace{I[y^i \neq h_t(x^i)]}_{= \begin{cases} 1 & \text{if } y^i \neq h_t(x^i) \\ 0 & \text{otherwise} \end{cases}}$$

						
1/16	1/4	1/16	1/16	1/4	1/16	1/4
✓	✓	✓	✗	✗	✓	✓

$$\varepsilon_t = \frac{1}{4} + \frac{1}{16} = \frac{5}{16}$$

*If classified correctly
the weight of that ex
will be multiplied by 0
and thus not included*

- ε_t is the weight of all misclassified examples added
 - The error rate is computed over original examples, not the re-sampled examples
- If a weak classifier is better than random, then $\varepsilon_t < 1/2$

AdaBoost: Step 3

[3] Compute weight α_t of classifier h_t

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_t}{\varepsilon_t} \right)$$

- In example from previous slide error was 5/16

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - 5/16}{5/16} \right) = \frac{1}{2} \ln \left(\frac{0.6875}{0.3125} \right) = 0.394 \approx 0.4$$


- Recall that $\varepsilon_t < 1/2$ if weak classifier is better than random
- The smaller is ε_t , the larger is α_t , and thus the more importance (weight) classifier $h_t(x)$

AdaBoost: Step 4

[4] For each x^i , update weights for next iteration


$$D_{t+1}(x^i) = D_t(x^i) \cdot \exp(\alpha_t \cdot I[y^i \neq h_t(x^i)])$$

- From previous slide $\alpha_t = 0.4$



- Weight of misclassified examples is increased

AdaBoost: Step 4



- First example (face) increased to 0.093 from 0.062 (which was $\frac{1}{16}$)
- Second example (face) increased to 0.373 from 0.25

AdaBoost: Step 5


[5] Normalize $D(x^i)$ so that $\sum D(x^i) = 1$

- *Then start over!*

AdaBoost Example


- Initialization: all examples have equal weights

from "A Tutorial on Boosting" by Yoav Freund and Rob Schapire




AdaBoost Example

ROUND 1




$$h_1(x) = \text{sign}(3 - x_1)$$

$$\begin{aligned}\epsilon_1 &= 0.30 \\ \alpha_1 &= 0.42\end{aligned}$$




AdaBoost Example

ROUND 2




AdaBoost Example

ROUND 3



$$h_3(\mathbf{x}) = \text{sign}(x_2 - 4)$$

AdaBoost Example



$$\text{sign}(0.42\text{sign}(3 - x_1) + 0.65\text{sign}(7 - x_1) + 0.92\text{sign}(x_2 - 4))$$

AdaBoost Advantages


- Can construct arbitrarily complex decision regions
- Fast
- Simple
- Has only one parameter to tune, T
- Flexible: can be combined with any classifier
- Provably effective (assuming weak learner)
 - Shift in mind set: goal now is merely to find hypotheses that are better than random guessing!
- While overfitting could occur, it would need a lot of iterations (which isn't always necessary for final classifier)

Caveats

- AdaBoost can fail if
 - If weak hypothesis (weak learners) are too complex
 - Overfitting
 - If weak hypothesis (weak learners) are too weak
 - Underfitting
- Empirically, AdaBoost seems especially susceptible to noise
 - Data with wrong labels

Test Error // AdaBoost

- Think about overfitting/underfitting
- Occam's Razor



Sample Snippets of Code

```
from sklearn.ensemble import AdaBoostClassifier #For Classification
from sklearn.ensemble import AdaBoostRegressor #For Regression
from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier()
clf = AdaBoostClassifier(n_estimators=100, base_estimator=dt, learning_rate=1)
# Decision tree was used as a base estimator, you can use any
# ML learner as base estimator if it accepts sample weight
clf.fit(x_train, y_train)
```

You can tune the parameters to optimize the performance of algorithms Some key parameters for tuning mentioned below:

n_estimators: It controls the number of weak learners.

learning_rate: Controls the contribution of weak learners in the final combination. There is a trade-off between learning_rate and n_estimators.

base_estimators: It helps to specify different ML algorithm.

You can also tune the parameters of base learners to optimize its performance.



Additional Material



More on Bagging

- The perturbation in the training set due to the bootstrap resampling causes different hypotheses to be built, particularly if the classifier is unstable
 - A classifier is said to be unstable if a small change in the training data (e.g., order of presentation of examples) can lead to a radically different hypothesis. This is the case of decision trees and, arguably, neural networks
- Bagging can be expected to improve accuracy if the induced classifiers are uncorrelated
 - In some cases, such as k Nearest Neighbors, bagging has been shown to degrade performance as compared to individual classifiers as a result of an effectively smaller training set
- A related approach to bagging is “cross-validated committees”, in which the component classifiers are built on different partitions of the training set obtained through k-fold cross-validation

Another Description of AdaBoost

AdaBoost operates as follows

- At iteration n , boosting provides the weak learner with a distribution D_n over the training set, where $D_n(i)$ represents the probability of selecting the i -th example
 - The initial distribution is uniform: $D_1(i) = 1/N$. Thus, all examples are equally likely to be selected for the first component
- The weak learner subsamples the training set according to D_n and generates a trained model or hypothesis H_n
- The error rate of H_n is measured with respect to the distribution D_n
- A new distribution D_{n+1} is produced by decreasing the probability of those examples that were correctly classified, and increasing the probability of the misclassified examples
- The process is repeated T times, and a final hypothesis is obtained by weighting the votes of individual hypotheses $\{h_1, h_2, \dots, h_T\}$ according to their performance

Note

- The strength of AdaBoost derives from the adaptive re-sampling of examples, not from the final weighted combination
 - To prove this point Breiman developed a variant of AdaBoost, known as 'arc-x4', in which the ensemble voting is unweighted [Breiman, 1996]: his results show that AdaBoost (referred to as 'arc-fs') and 'arc-x4' have similar performance [Bauer and Kohavi, 1999]

The Bias and Variance Decomposition

The effectiveness of Bagging and Boosting has been explained in terms of the bias-variance decomposition of classification error

- The expected error of a learning algorithm can be decomposed into
 - A **bias** term that measures how closely the average classifier produced by the learning algorithm matches the target function
 - A **variance** term that measures how much the learning algorithm's predictions fluctuate for different training sets (of the same size)
 - An **intrinsic target noise**, which is the minimum error that can be achieved: that of the Bayes optimal classifier
- Following this line of reasoning, Breiman has suggested that both Bagging and Boosting reduce errors by reducing the variance term
- Along the same lines, Freund and Schapire have argued that Boosting also attempts to reduce the bias term since it focuses on misclassified samples
 - Work by Bauer and Kohavi, however, seems to indicate that Bagging can also reduce the bias term

References:

- ~Ensemble Learning ~ R. Gutierrez-Osuna ~ Pattern Analysis ~ TAMU
- ~Analytics Vidhya ~ Quick Introduction to Boosting Algorithms in Machine Learning ~ S.Ray
- ~AdaBoost Tutorial ~ C.McCormick
- ~Boosting ~ Machine Learning ~ O.Veksler
- ~Boosting (AdaBoost Algorithm) ~ E.Emmer
- ~Excellent paper on AdaBoost written by one of the original authors of the algorithm, *Robert Schapire*: <http://rob.schapire.net/papers/explaining-adaboost.pdf>
- ~A Tutorial on Boosting ~ Y.Freund and R.Schapire ~ di.unipi.it