



# **Deep Learning**

## **Convolutional Neural Networks**

***CS 6316 – Machine Learning***  
***Fall 2017***

Adapted from material from Prof. Qi and Prof. Yann LeCun et al. (Facebook AI research / NYU)

# Deep Learning

(Quote from Yann LeCun et al. – Nature, Vol. 521, 2015)

*Machine Learning technology powers many aspects of modern society: from web searches to content filtering on social networks to recommendations on e-commerce websites, and it is increasingly present in consumer products such as cameras and smartphones.*

# Deep Learning

(Quote from Yann LeCun et al. – Nature, Vol. 521, 2015)

*Machine Learning systems are used to identify objects in images, transcribe speech into text, match news items, posts or products with users' interests, and select relevant results of search.*

*Increasingly, these applications make use of a class of techniques called **deep learning**.*

# Convolutional Neural Networks and Deep Learning

- Convolutional neural networks (CNN) are a type of artificial neural network (ANN) that includes both **fully-connected layers** and **locally-connected layers** known as **convolutional layers**

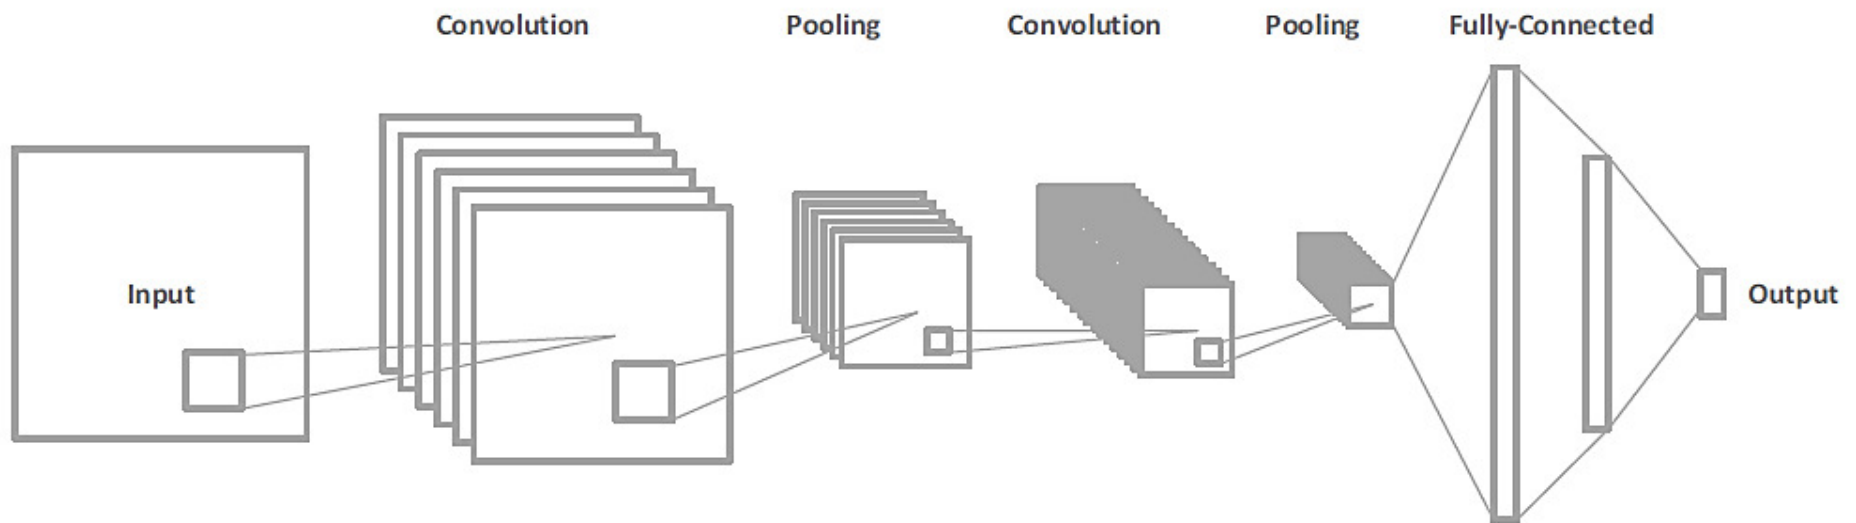


Figure 1: Convolutional neural network composed of convolution, pooling, and fully-connected layers

# Convolutional Neural Networks and Deep Learning

- In large (“*deep*”) convolutional networks, it is common to see other types of layers such as pooling layers, activation layers, and batch normalization layers

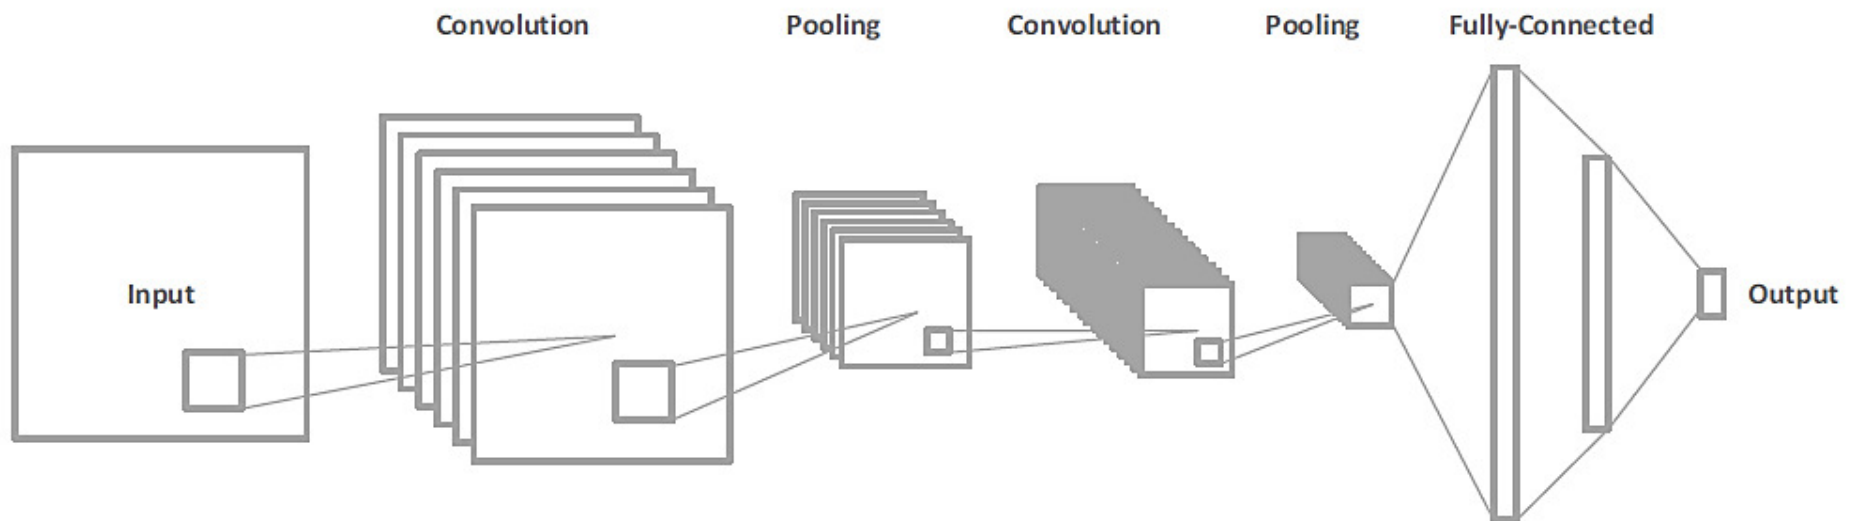


Figure 1: Convolutional neural network composed of convolution, pooling, and fully-connected layers



# Convolutional Neural Networks and Deep Learning

- Here we see the architecture of a simple convolutional neural network consisting of **two convolutional layers**, **two pooling layers**, and **three fully connected layers**:

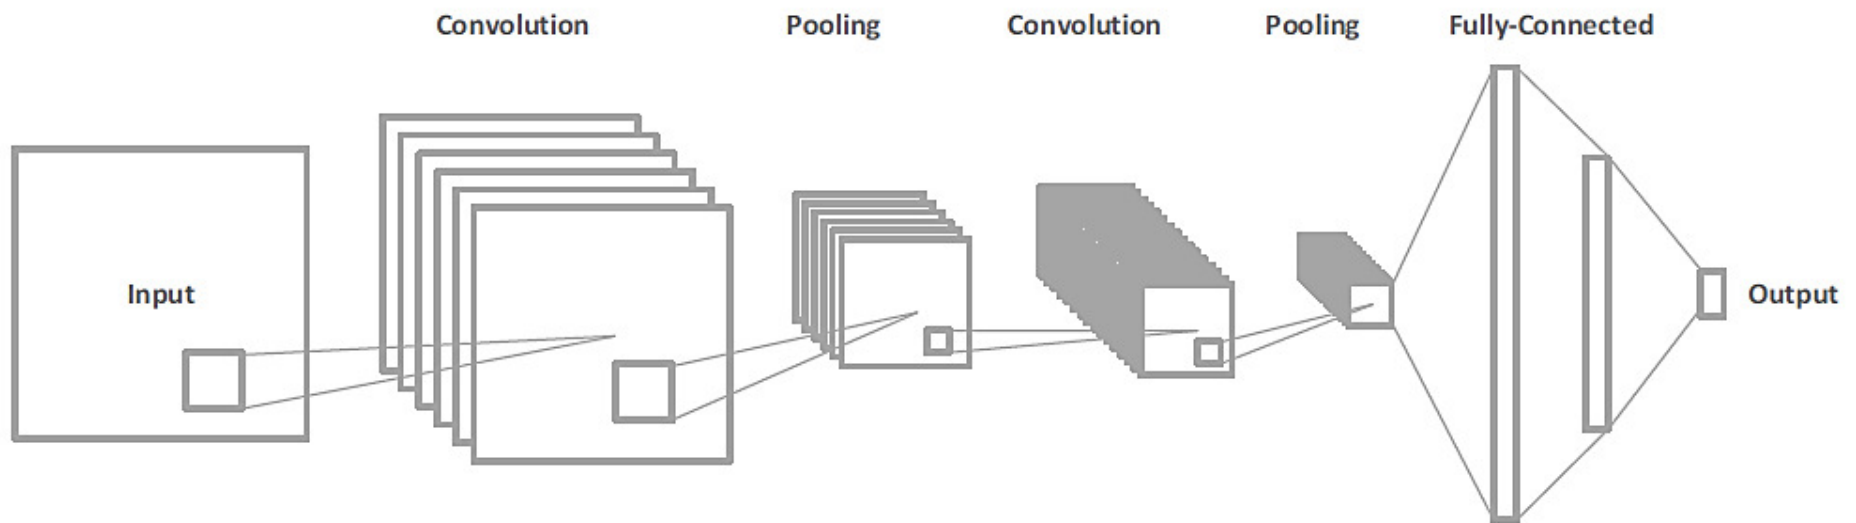


Figure 1: Convolutional neural network composed of convolution, pooling, and fully-connected layers

# Data: Multiple Arrays

- Convolutional Neural Networks are designed to process data that come in the form of multiple arrays, for example a color image composed of three 2D arrays containing pixel intensities in the three color channels

# Data: Multiple Arrays

- Many data modalities are in the form of **multiple arrays**:
  - 1D for signals and sequences, including language
  - 2D for images or audio spectrograms
  - 3D for video or volumetric images
- There are four key ideas behind CNNs that take advantage of the properties of natural signals:
  - Local connections
  - Shared weights
  - Pooling
  - Use of many layers



# Convolutional Neural Networks

- Training deep convolutional neural networks from scratch is **difficult** since training can require
  - extensive computational resources and
  - large amounts of training data
- If such resources are not available, a **pre-trained CNN** can be used to perform a new classification task by either
  - [A] fine-tuning the pre-trained network's weights using a new target dataset, or by
  - [B] using the pre-trained network's activations as automatic feature extractors.

# Convolutional Neural Networks

- Pre-trained networks for many popular CNN architectures such as AlexNet (A.Krizhevsky et al.), VGGNet (K.Simonyan and A Zisserman), and GoogLeNet (C.Szegedy et al. ) can be downloaded and adapted for different problems
- Many applications, but popular application is visual (image) recognition

What we see ...



What we see ...



# What we see ...



# Convolutional Neural Networks (aka CNNs and ConvNets)

- **The problem:**  
*semantic gap*
- Images are represented as 3D arrays of numbers, with integers between  $[0, 255]$
- E.g.  $300 \times 300 \times 3$  (3 for e color channels RGB)



08	02	22	97	38	15	00	40	00	75	04	05	07	78	52	22	50	77	51	88
49	49	99	40	17	81	18	57	60	87	17	40	98	43	69	44	04	56	62	00
81	49	31	73	55	79	14	29	93	71	40	67	53	55	30	03	49	13	36	45
52	70	95	23	04	40	11	42	65	84	68	56	01	32	56	71	37	02	36	91
22	31	16	71	51	67	73	59	41	92	36	54	22	40	40	28	66	33	13	80
24	47	38	63	99	03	45	02	44	75	33	53	78	36	84	20	35	17	12	50
52	98	81	28	64	23	67	10	26	38	40	47	59	54	70	66	18	38	64	70
67	26	20	68	02	62	12	20	95	63	94	39	63	08	40	91	66	49	94	21
24	55	58	05	66	73	99	26	97	17	78	78	96	83	14	88	34	89	63	72
21	36	23	09	75	00	76	44	20	45	35	14	00	61	33	97	34	31	33	95
78	17	53	28	22	75	31	67	15	94	03	80	04	62	16	14	09	53	56	92
16	39	05	42	94	35	31	47	55	58	88	24	00	17	54	24	34	29	85	57
86	56	00	48	35	71	89	07	05	44	44	37	44	60	21	58	51	54	17	58
19	80	81	68	05	94	47	69	28	73	92	13	86	52	17	77	04	89	55	40
04	52	08	83	97	35	99	16	07	97	57	32	16	26	26	79	33	27	98	66
57	56	68	87	57	62	20	72	03	46	33	67	46	55	12	32	63	93	53	69
04	42	16	73	35	39	11	24	94	72	18	08	46	29	32	40	62	76	36	
20	49	36	41	72	30	23	88	34	57	83	69	82	47	59	85	74	04	36	16
20	73	35	29	78	31	90	01	74	31	49	71	48	88	81	16	23	57	05	54
01	70	34	71	83	51	54	69	16	92	33	48	61	43	52	01	89	59	67	48

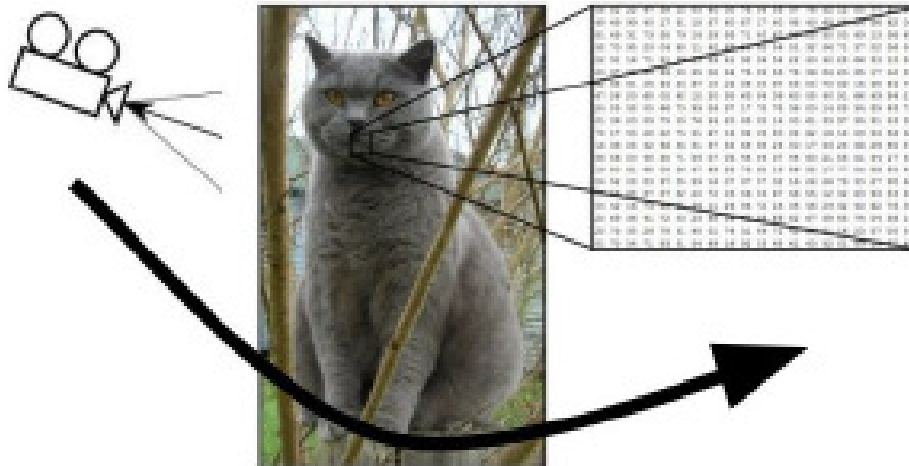
What the computer sees

**Challenges in  
Visual  
Recognition!**



# Challenges in Visual Recognition

- Camera pose
- Illumination



# Challenges in Visual Recognition

- Deformation
- Occlusion

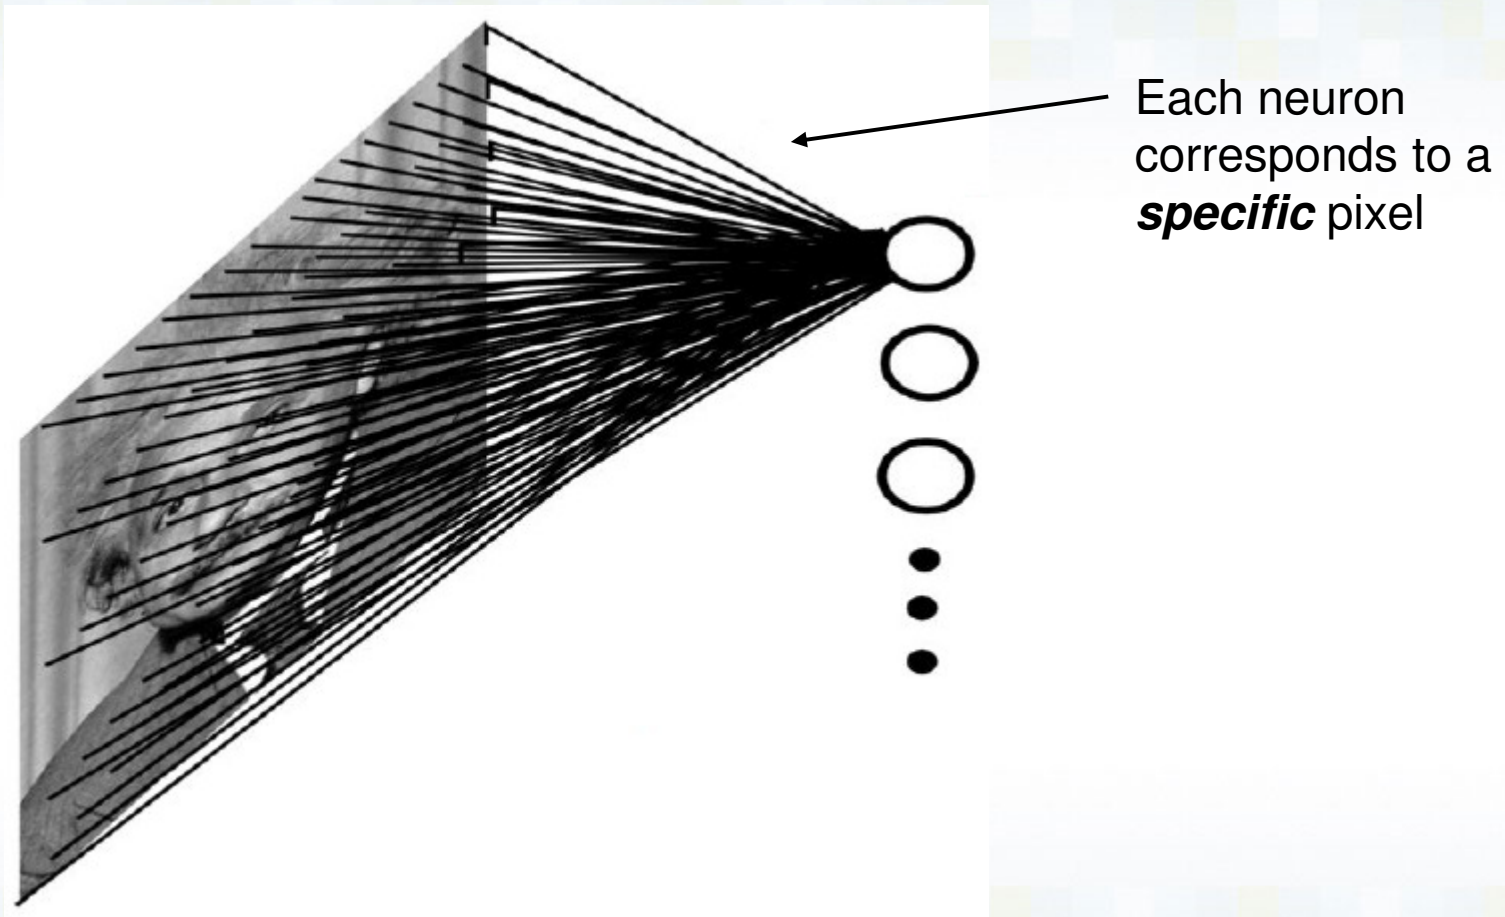


# Challenges in Visual Recognition

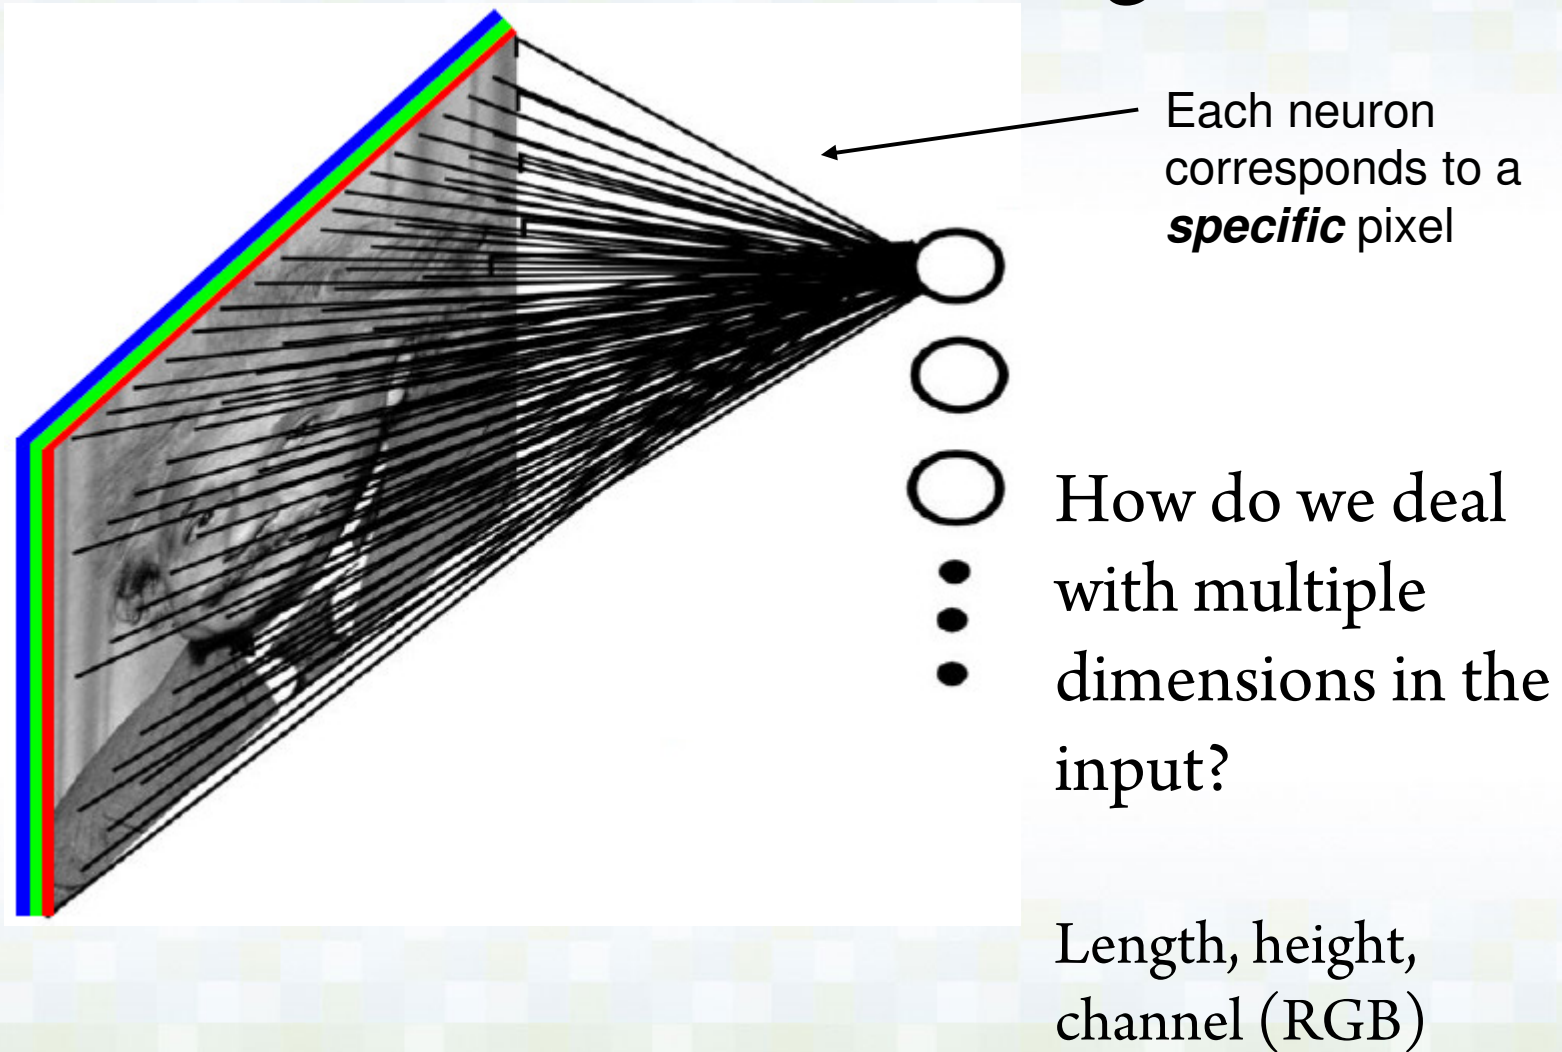
- Background clutter
- Intra-class variation



# Problems with “Fully Connected Networks” on Images

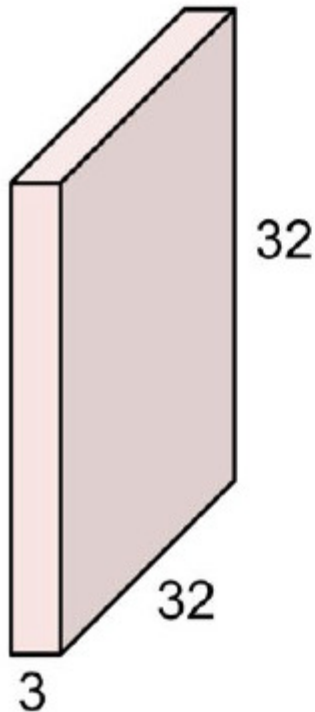


# Problems with “Fully Connected Networks” on Images



# Convolutional Layer

32x32x3 image



5x5x3 filter

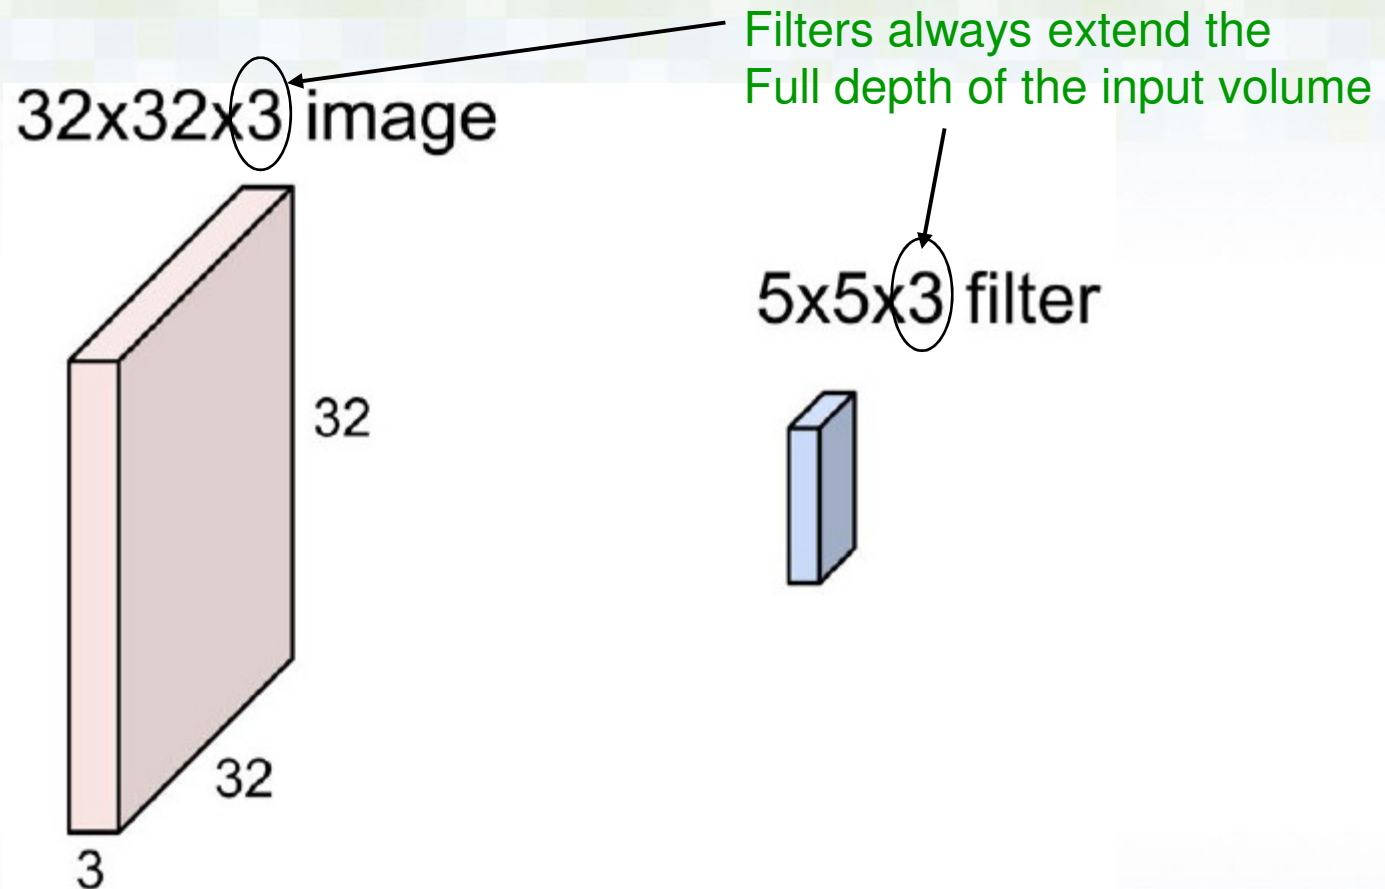


**Convolve** the filter with the image

i.e. “*slide over the image spatially, computing dot products*”



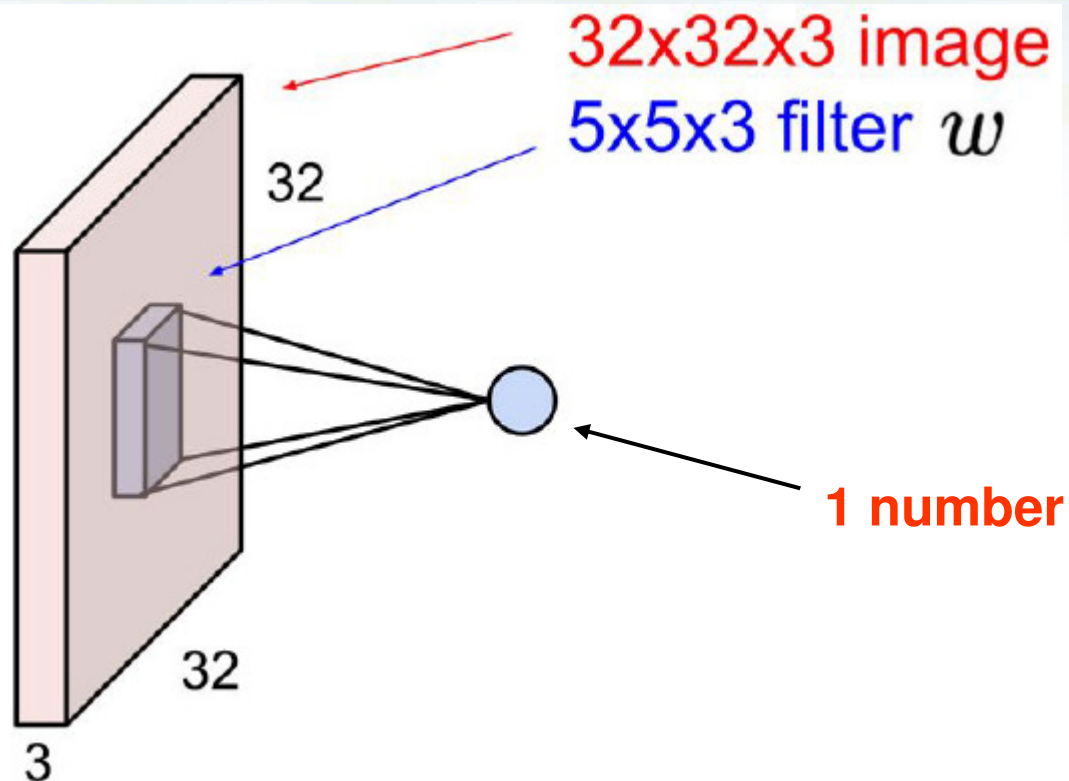
# Convolutional Layer



**Convolve** the filter with the image

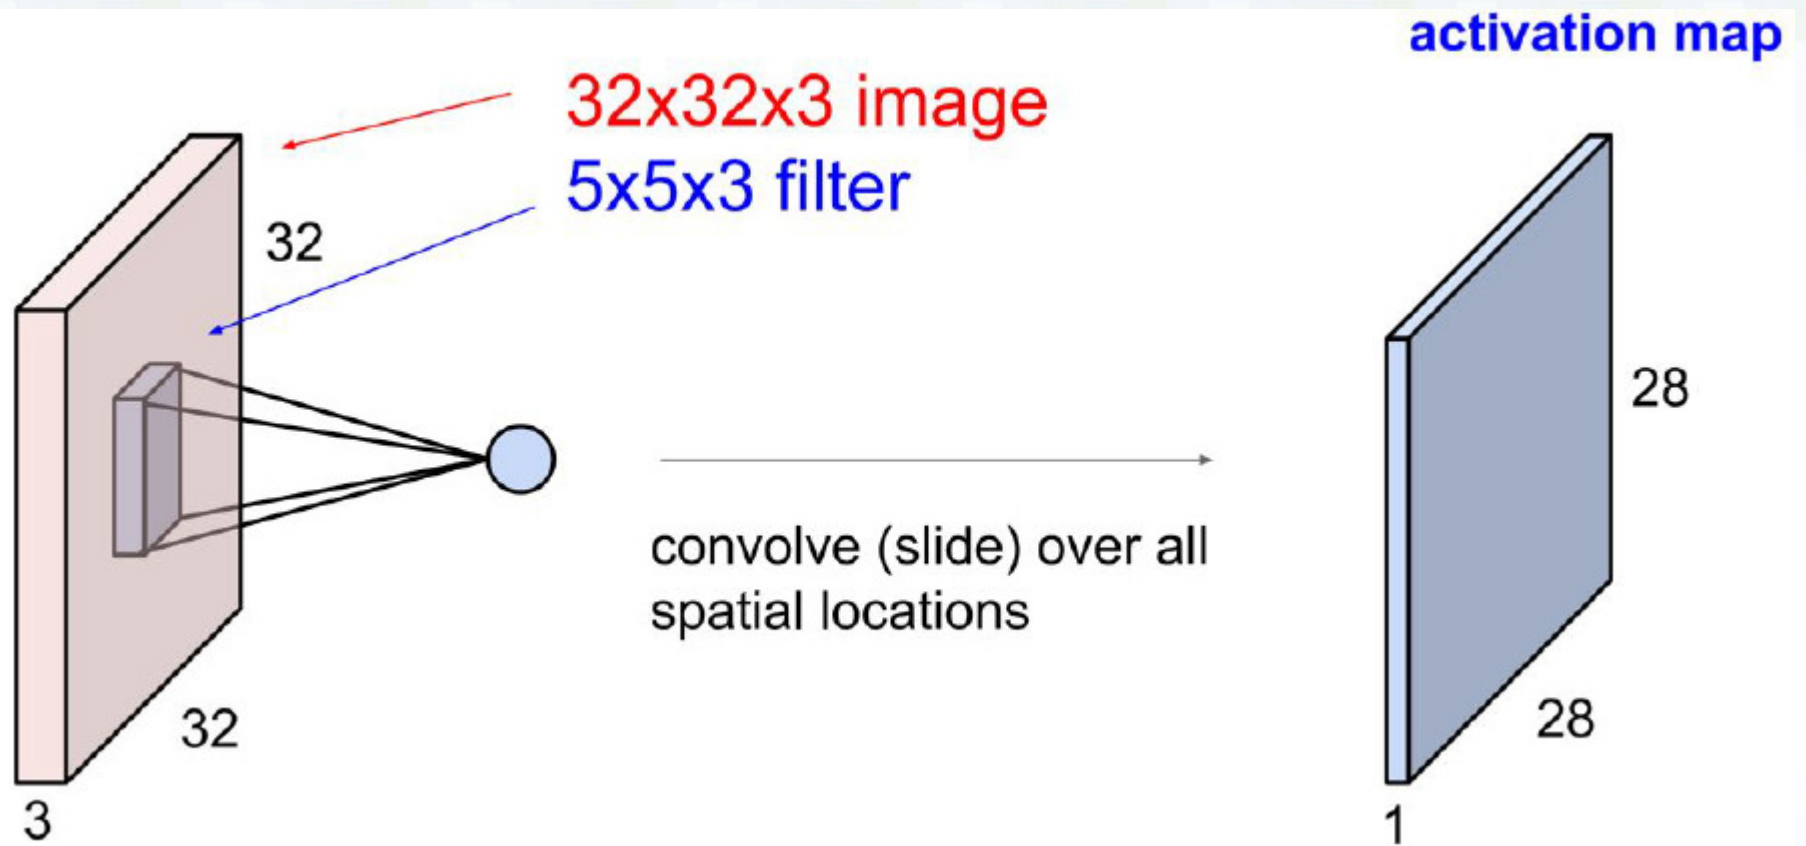
i.e. *“slide over the image spatially, computing dot products”*

# Convolutional Layer



- The result of taking a dot product between the filter and a small 5 x 5 x 3 chunk of the image (i.e.  $5*5*3 = 75$ -dimensional dot product+ bias)  $w^T x + b$

# Convolutional Layer



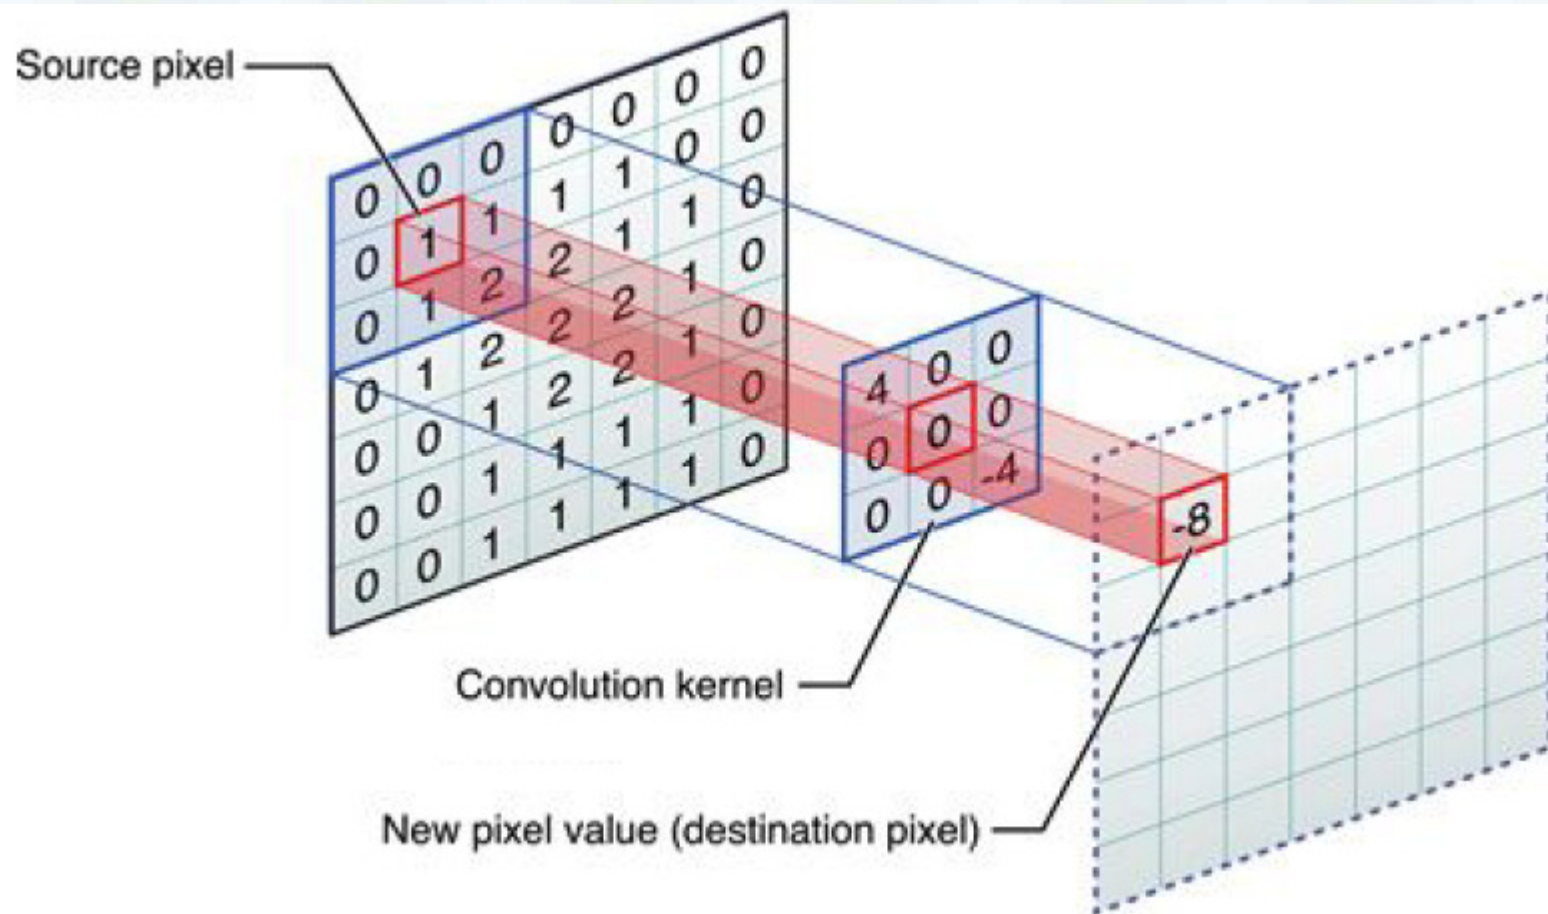
# Convolution

- We call the layer convolutional because it is related to convolution of two signals:

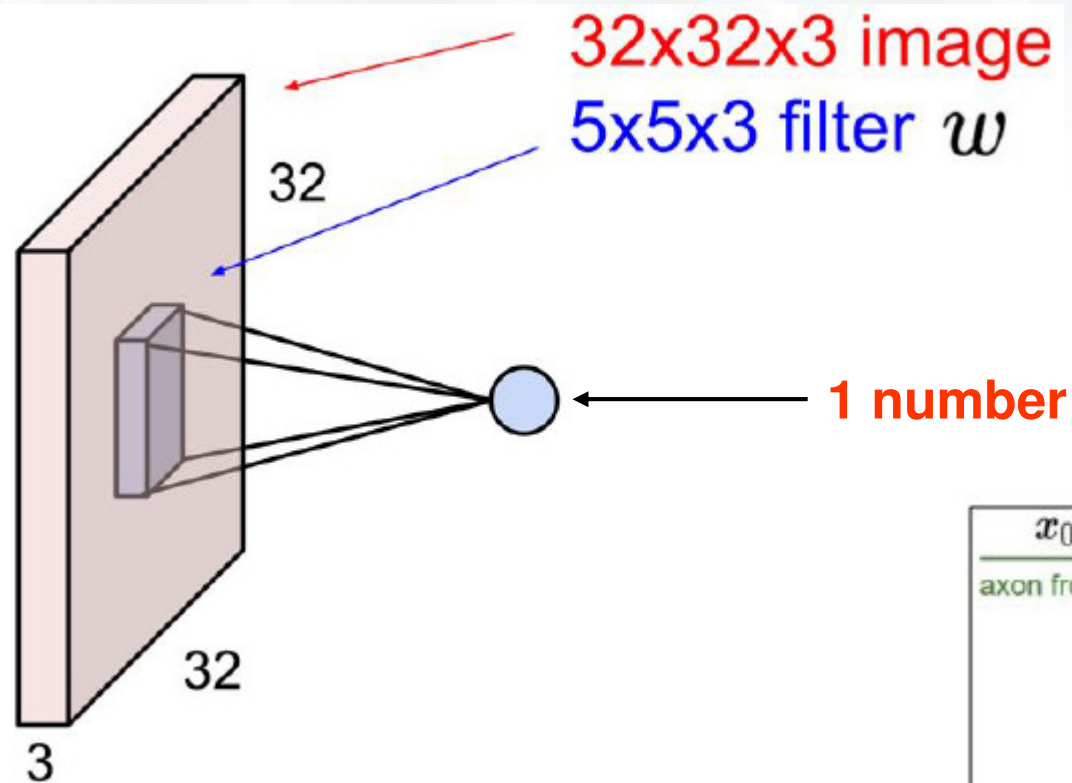
$$f[x,y] * g[x,y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1,n_2] \cdot g[x-n_1,y-n_2]$$

elementwise multiplication and sum of  
a filter and the signal (image)

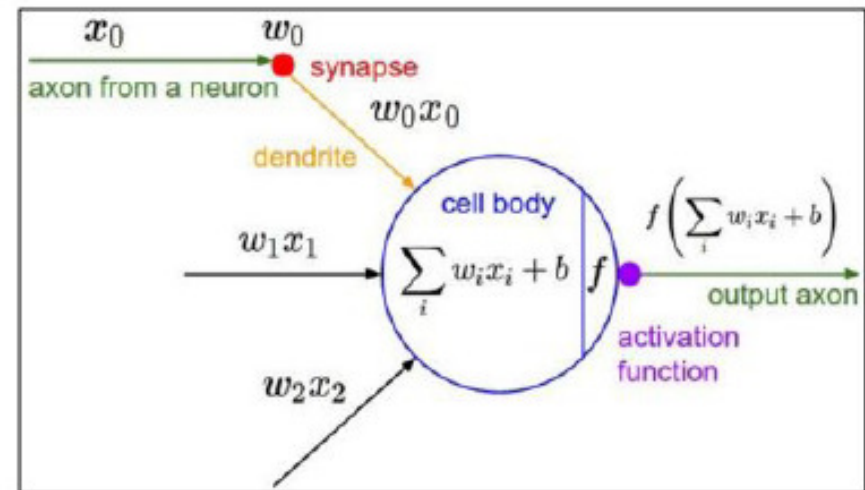
# Convolution



# Neuron View of Convolutional Layer

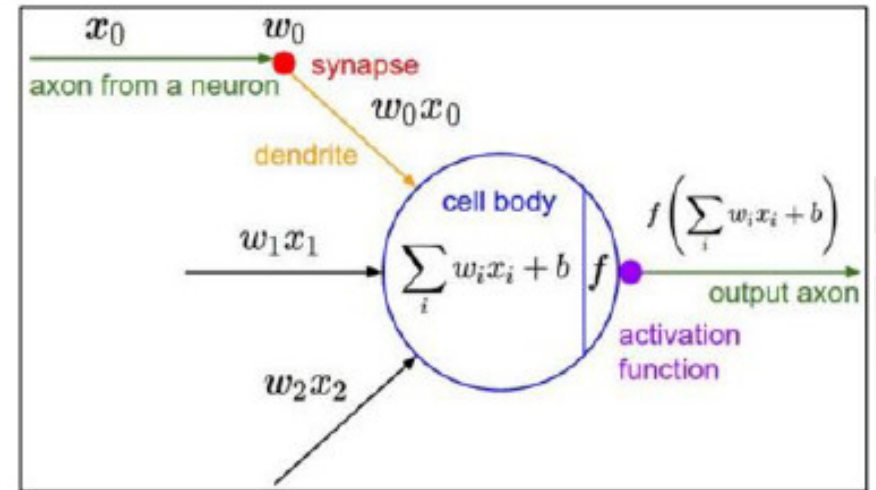


It is just a neuron with Local connectivity...

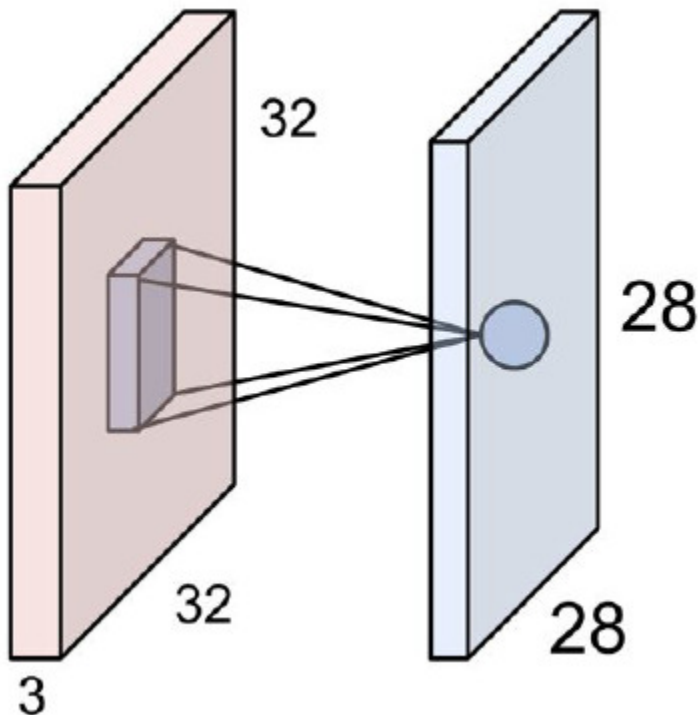




# Neuron View of Convolutional Layer



- An activation map is a 28 x 28 sheet of neuron outputs:
  1. Each is connected to a small region in the input
  2. All of them share parameters



“5x5 filter”  $\rightarrow$  “*5x5 receptive field for each neuron*”

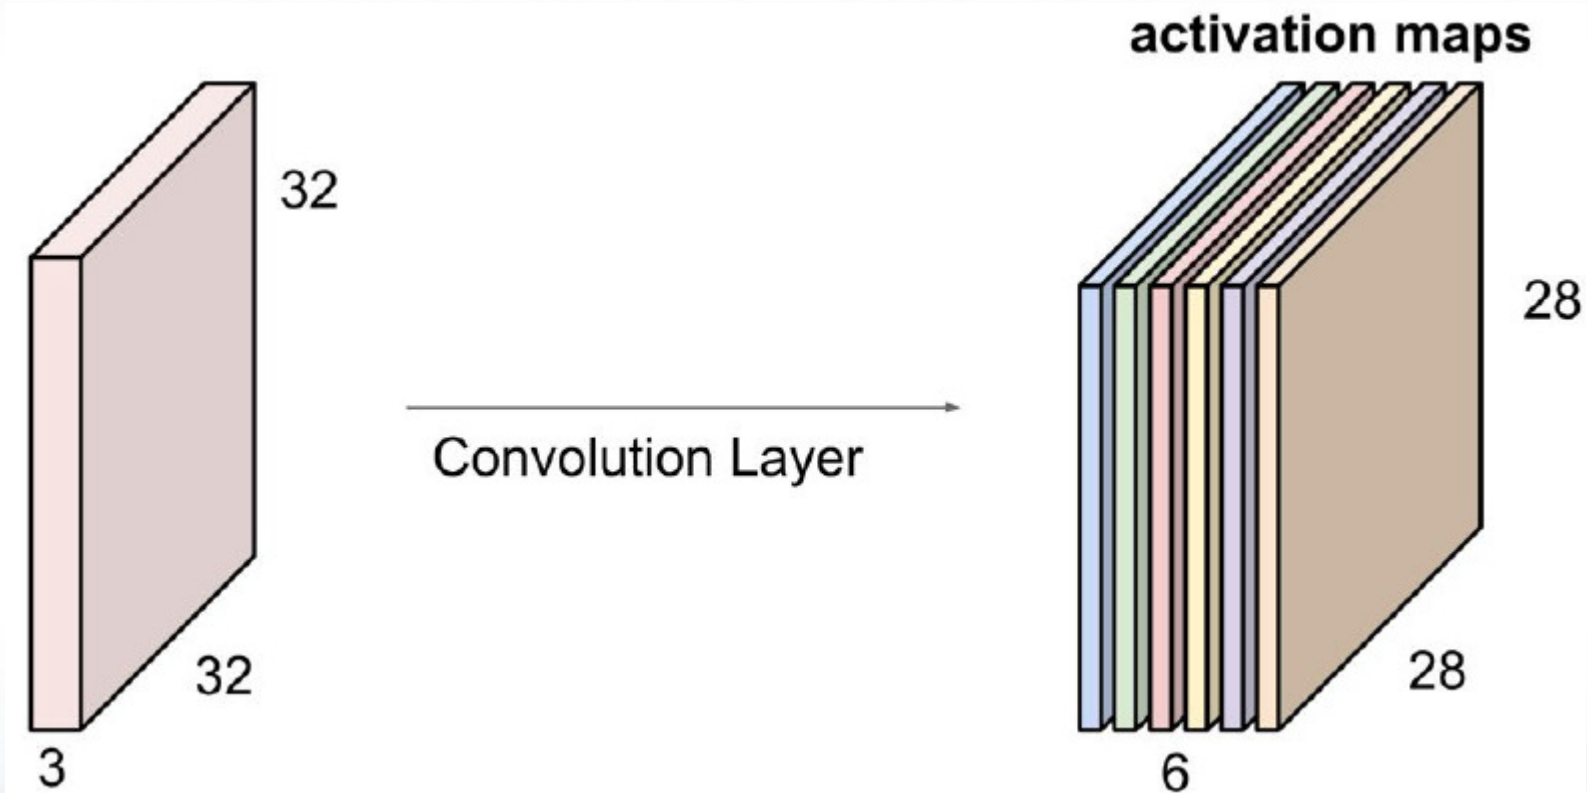
# Convolutional Layer

(consider a second, **GREEN** filter)



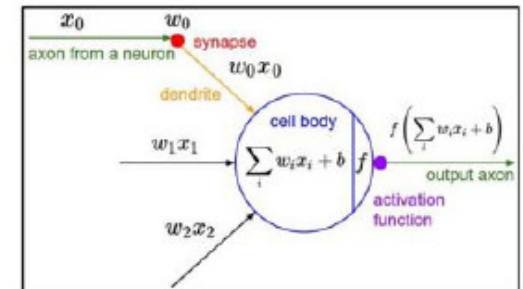
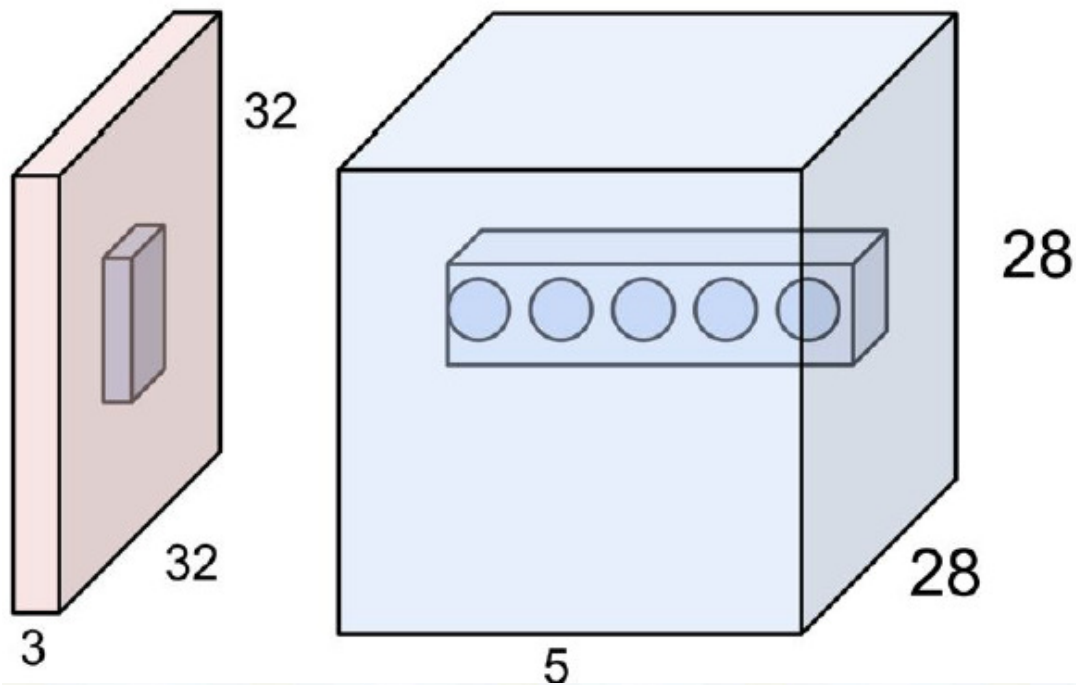
# Convolutional Layer

- For example, if we had **six (6) 5x5 filters**, we'll get 6 separate activation maps:



- We stack these up to get a “new image” of size **28 x 28 x 6!**

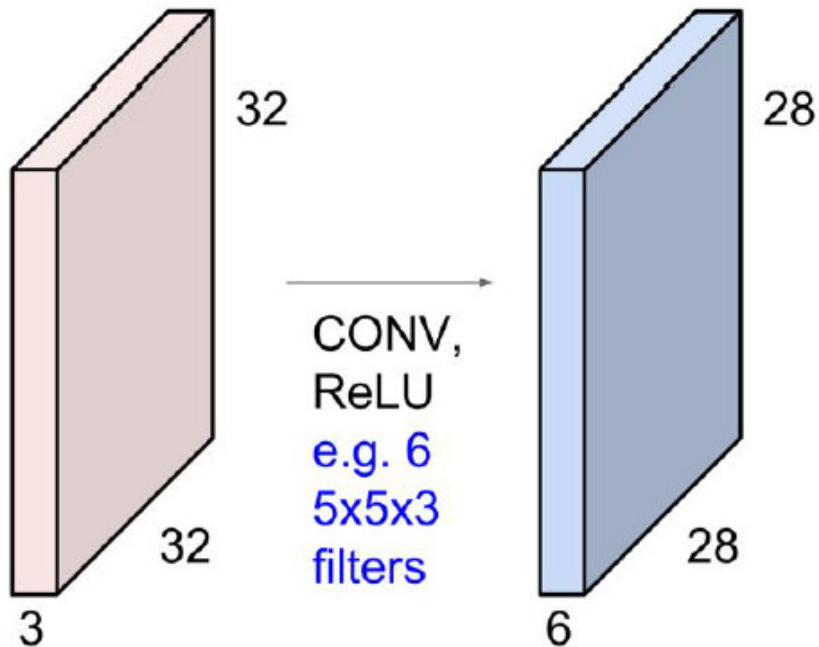
# Neuron View of Convolutional Layer



E.g. with 5 filters,  
CONV layer consists of  
neurons arranged in a 3D grid  
(28x28x5)

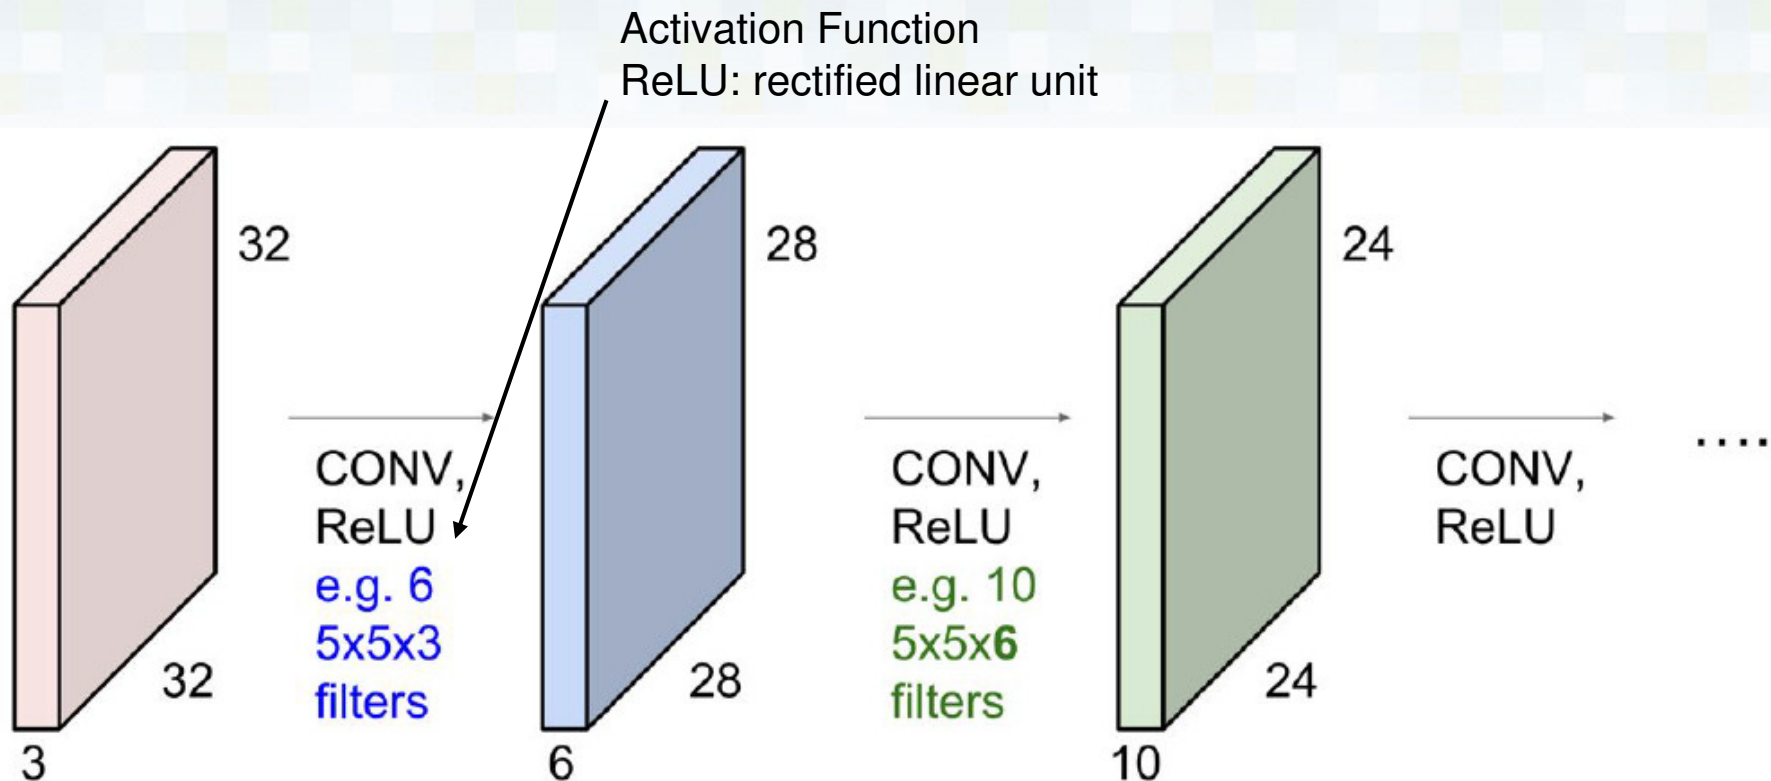
There will be 5 different  
neurons all looking at the same  
region in the input volume

# Convolutional Neural Networks



- Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions

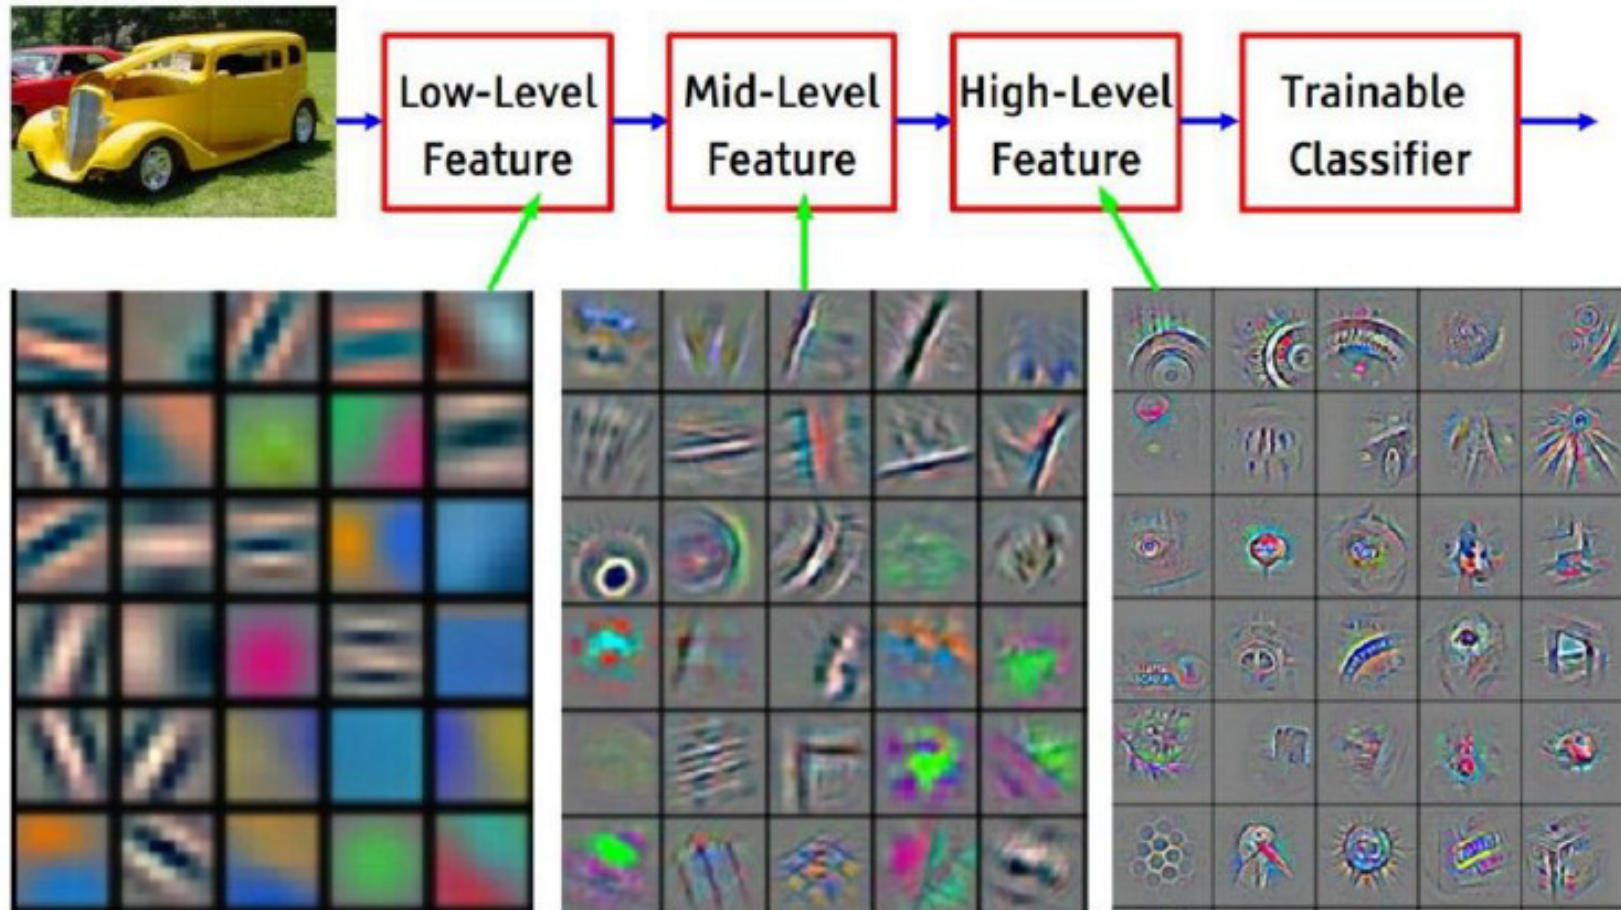
# Convolutional Neural Networks



- Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



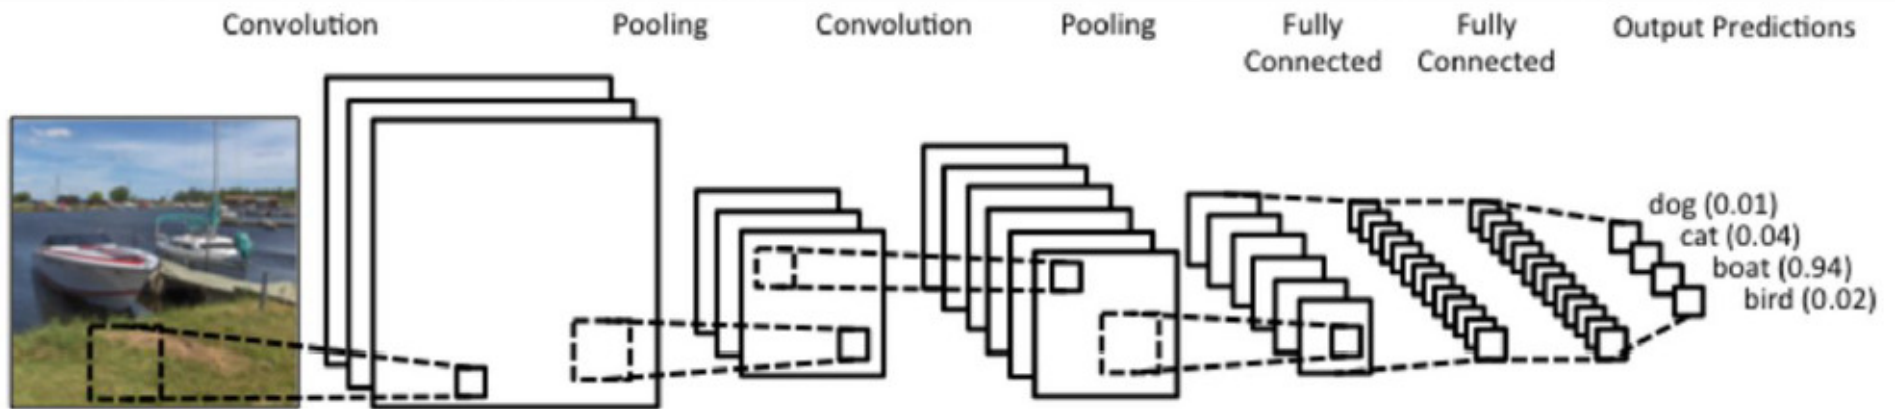
# Convolutional Neural Networks



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

[From recent Yann LeCun slides]

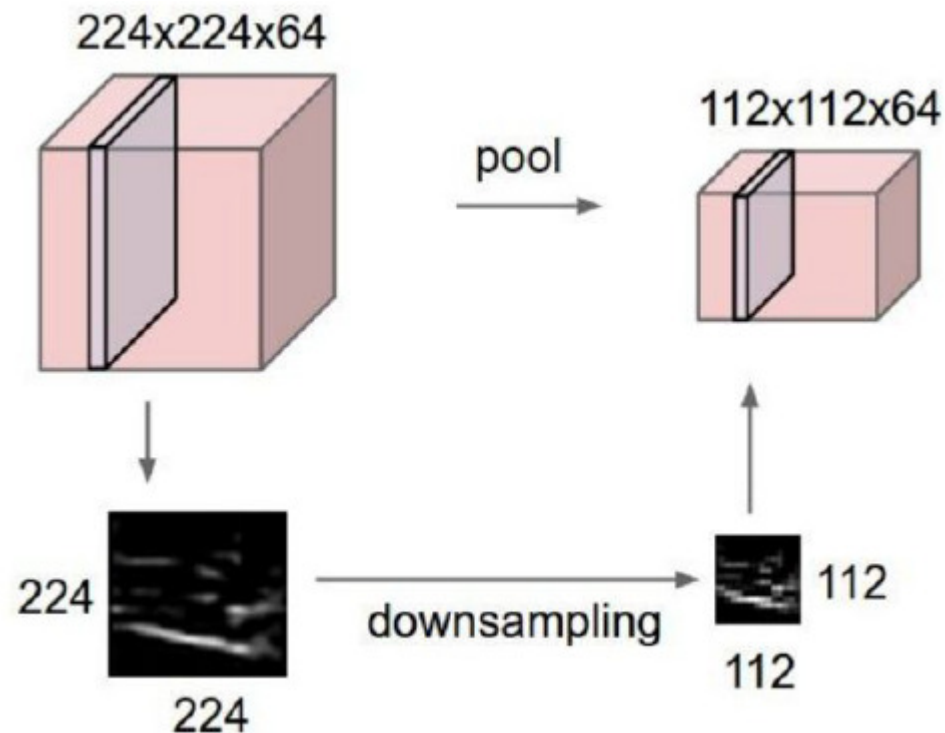
# Convolutional Neural Networks



<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>

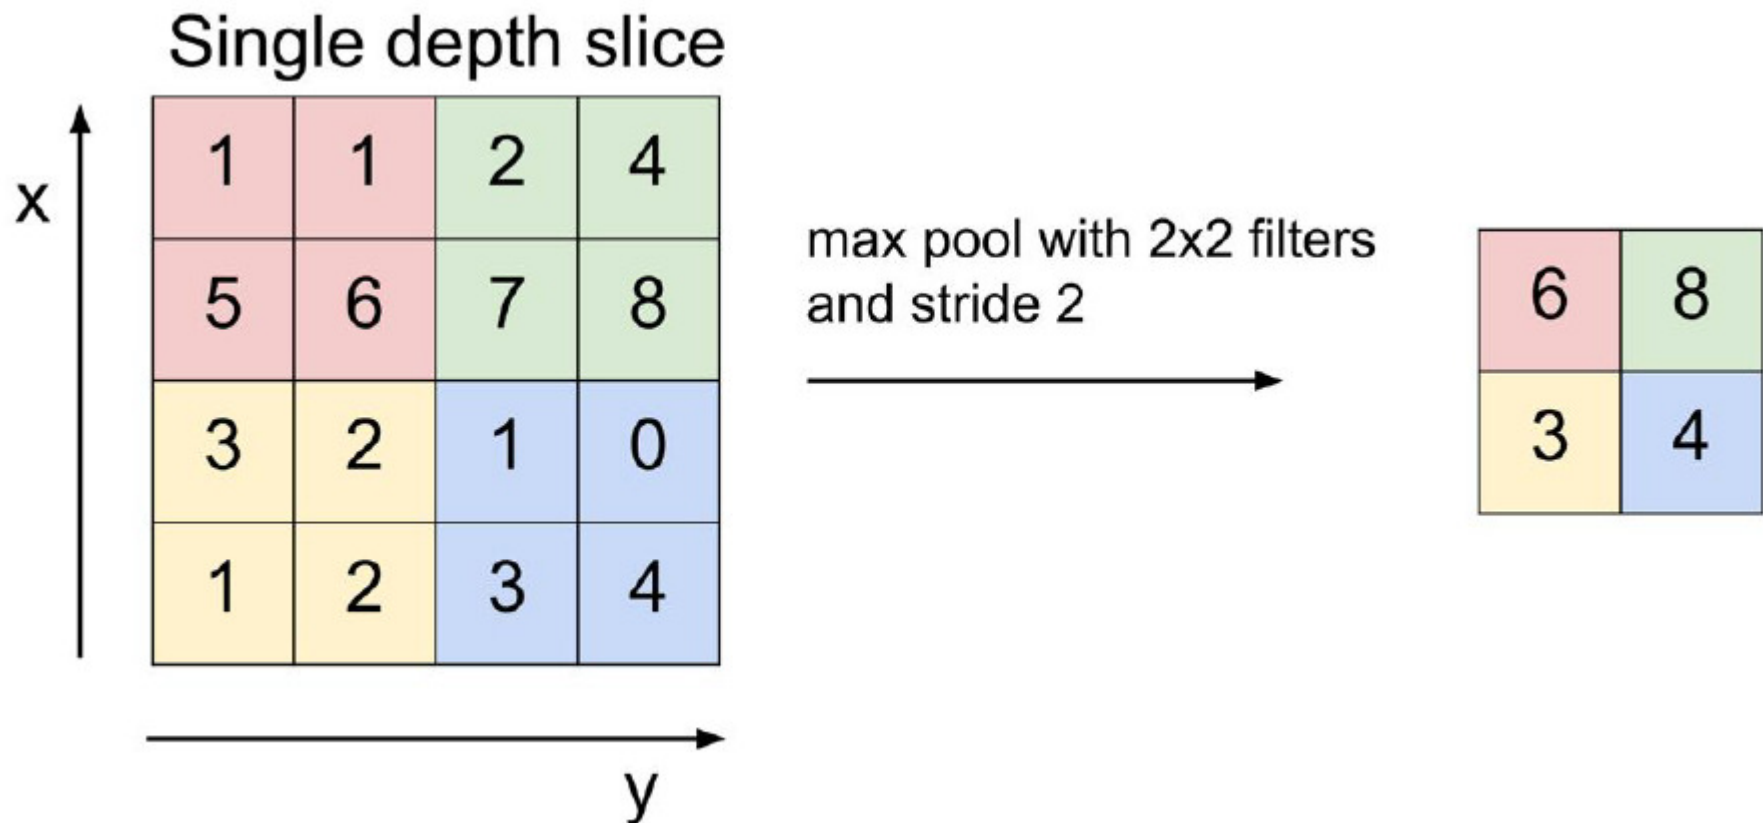
# Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently



# Pooling Layer (**Max Pooling** example)

- A typical pooling unit computes the maximum of a local patch of unites in one feature map (or in a few feature maps)



# Pooling Layer

- Although the role of the convolutional layer is to detect local conjunctions of features from the previous layer
- The role of the **pooling layer** is **to merge semantically similar features into one**
- Because the relative positions of the features forming a motif can **vary somewhat**, reliably detecting the motif can be done by **coarse-graining the position of each feature**

# Pooling Layer

The convolutional and pooling layers in ConvNets are directly inspired by the classic notions of **simple cells and complex cells in visual neuroscience**

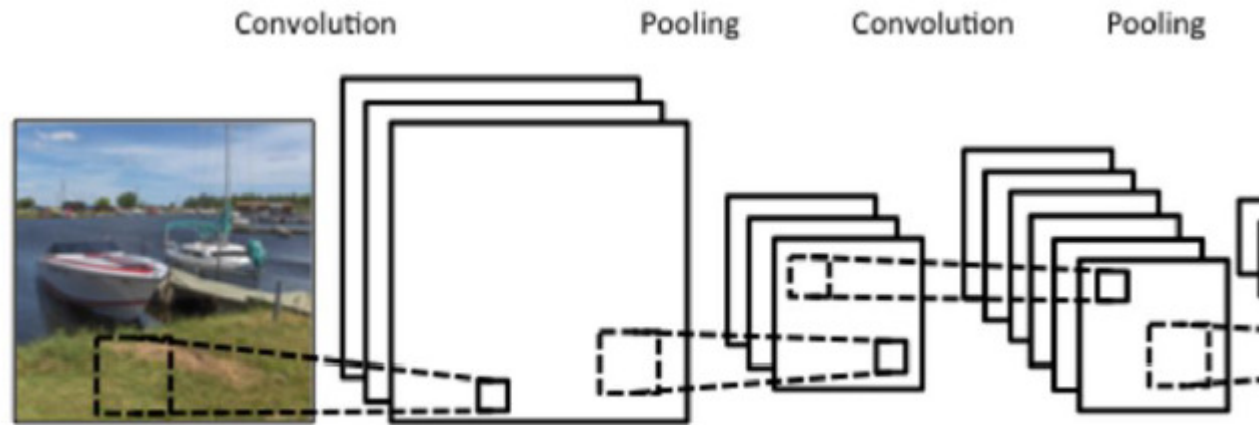


# Hierarchies

- Two or three stages of convolution and pooling are stacked, followed by more convolutional and fully-connected layers
- Deep neural networks exploit the property that many natural signals are **compositional hierarchies**, in which higher-level features are obtained by composing lower-level ones
- In images, local combinations of edges form motifs, motifs assemble into parts, and parts form objects

# Hierarchies and Pooling

- Neighboring pooling units take input from patches that are shifted by more than one row or column, thereby **reducing the dimension of the representation and creating an invariance to small shifts and distortions**

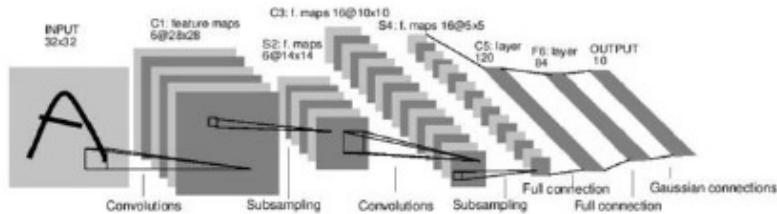


<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>

# History of CNNs

1998

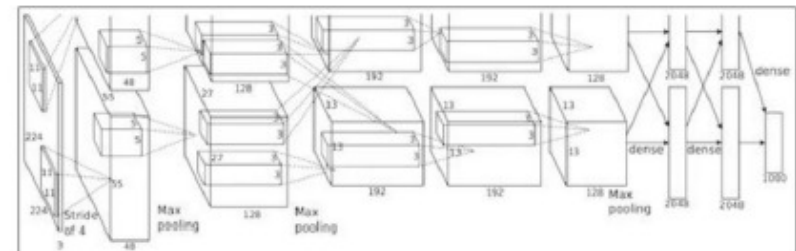
*Gradient-based learning applied to document recognition* [LeCun, Bottou, Bengio, Haffner]



LeNet-5

2012

*ImageNet Classification with Deep Convolutional Neural Networks* [Krizhevsky, Sutskever, Hinton, 2012]

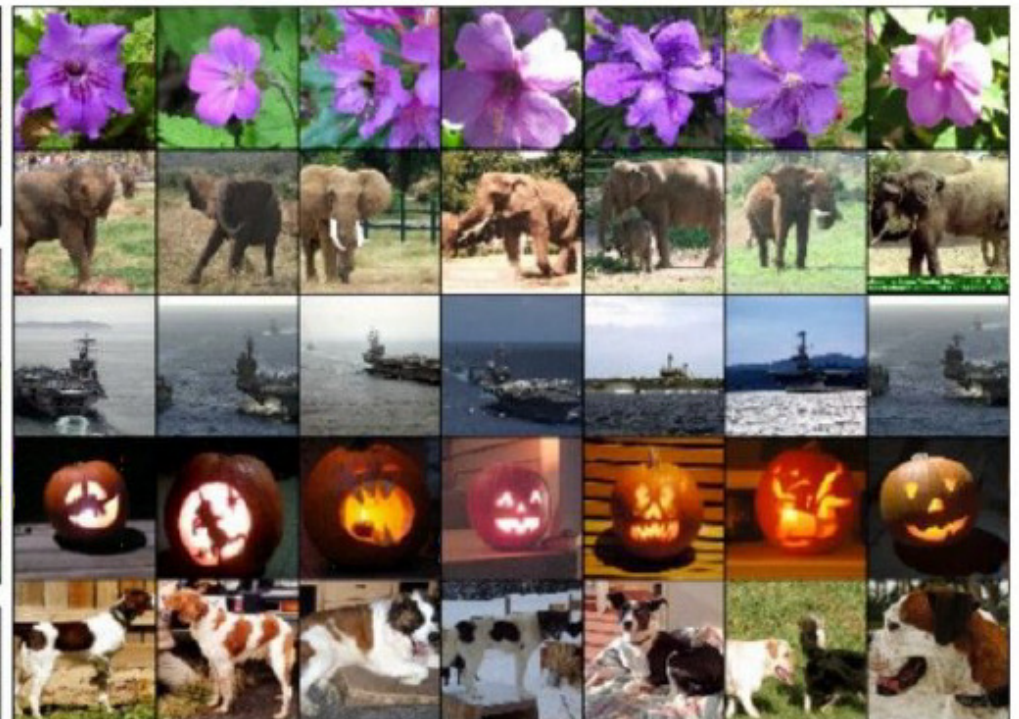
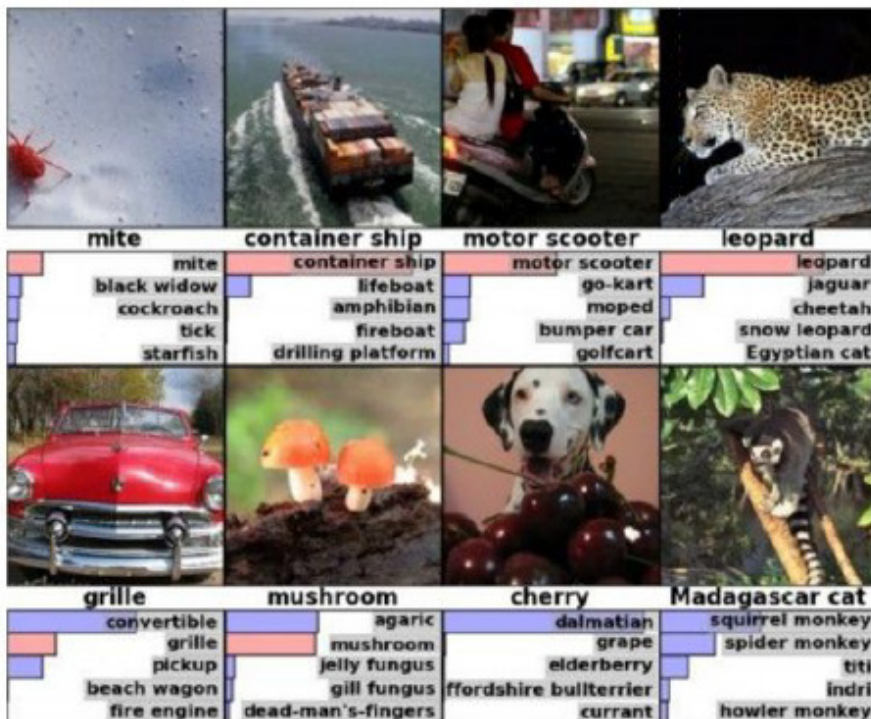


"AlexNet"

# Fast-forward to today: CNNs are everywhere

Classification

Retrieval

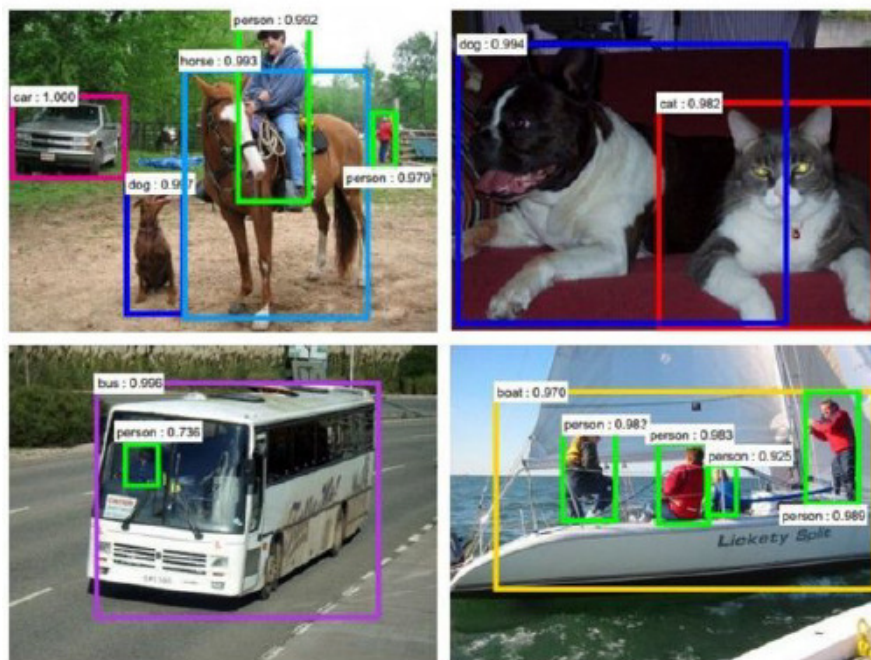


[Krizhevsky 2012]



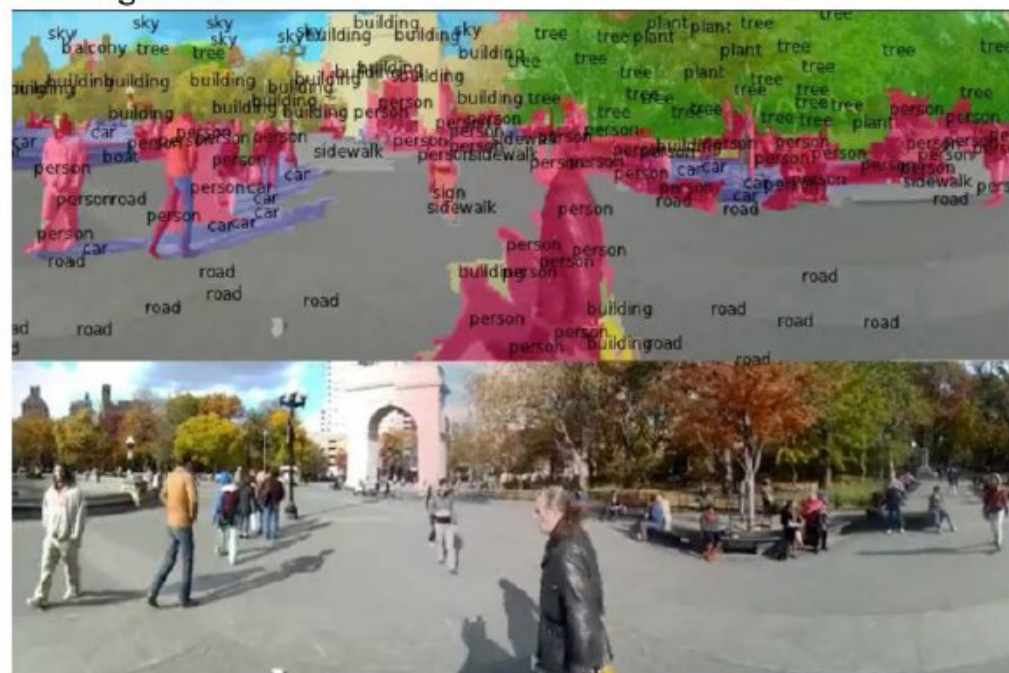
# Fast-forward to today: CNNs are everywhere

Detection



[Faster R-CNN: Ren, He, Girshick, Sun 2015]

Segmentation



[Farabet et al., 2012]

# ... everywhere



self-driving cars



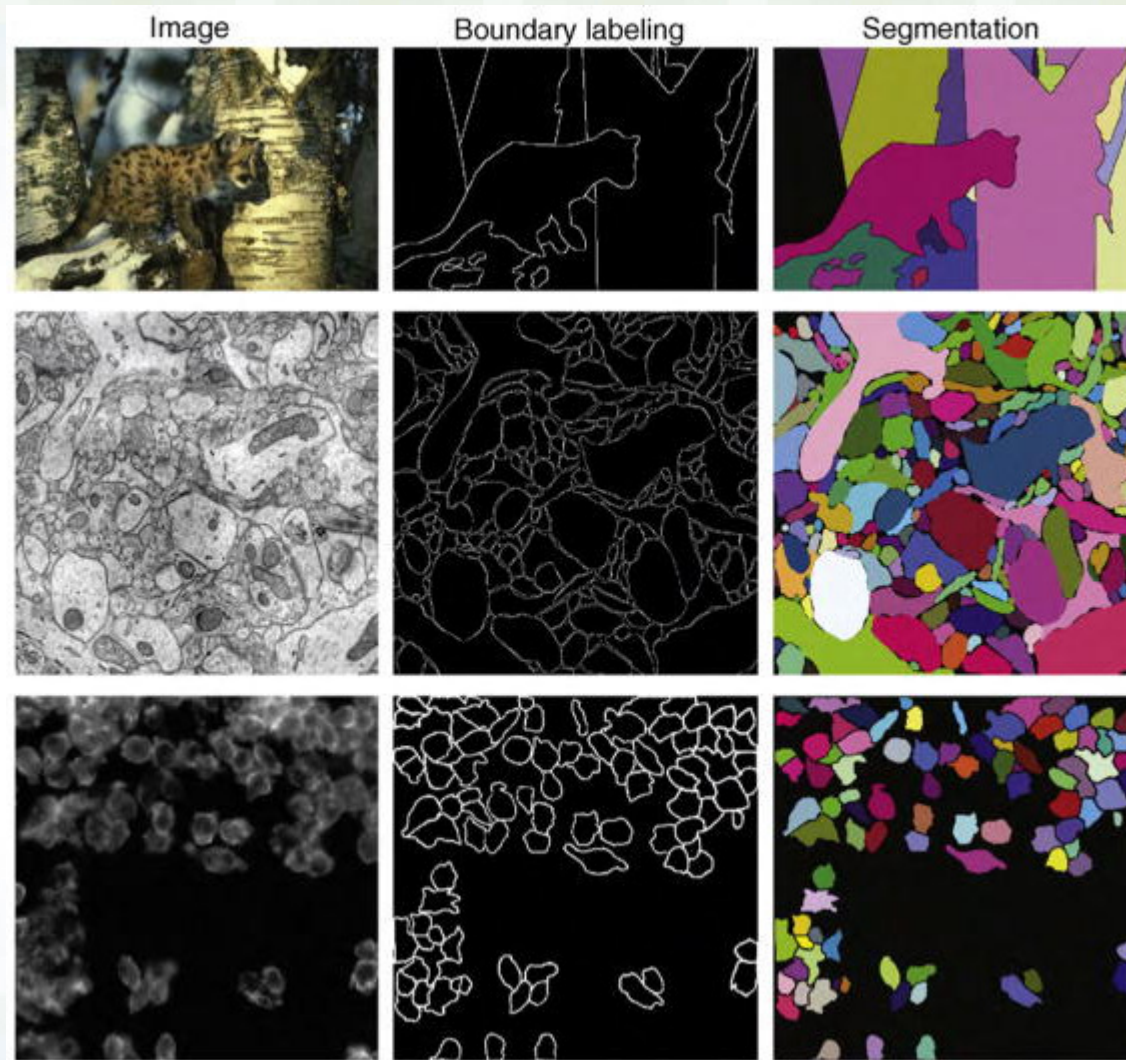
NVIDIA Tegra X1



# ...everywhere

## Segmentation

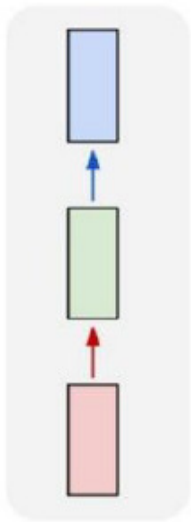
[Turaga et al.,  
2010]



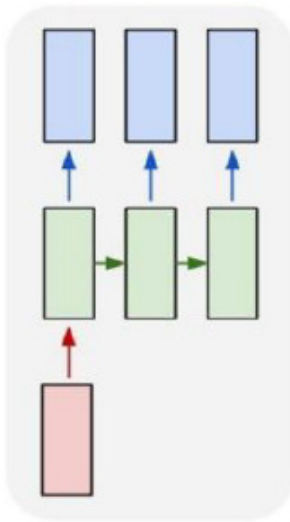
# Recurrent Neural Networks (*Brief*)

## Standard “Feed-Forward” Neural Network

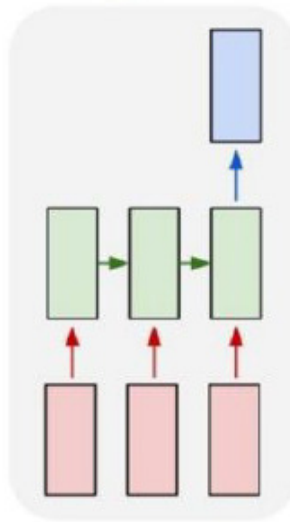
one to one



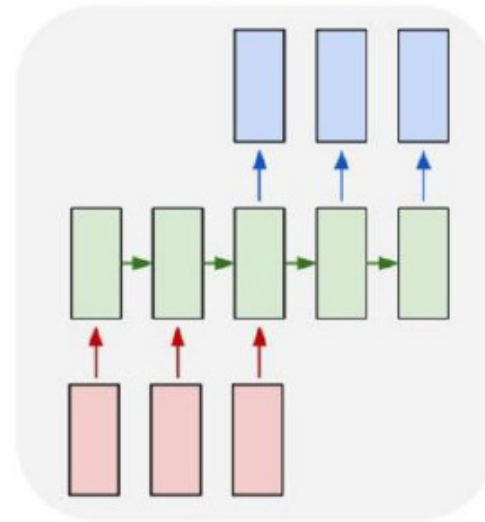
one to many



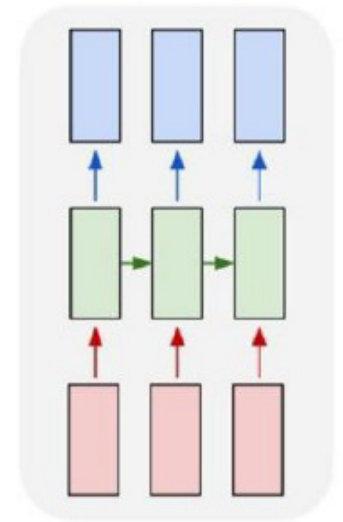
many to one



many to many



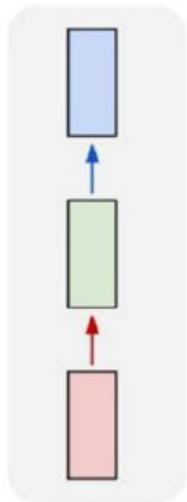
many to many



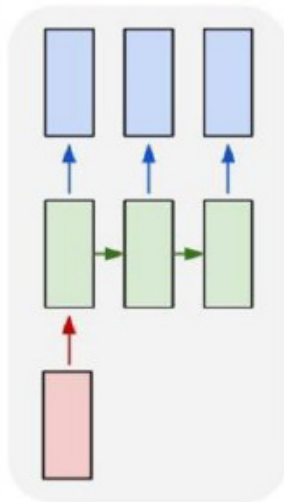
# Recurrent Neural Networks (RNNs)

RNNs can handle

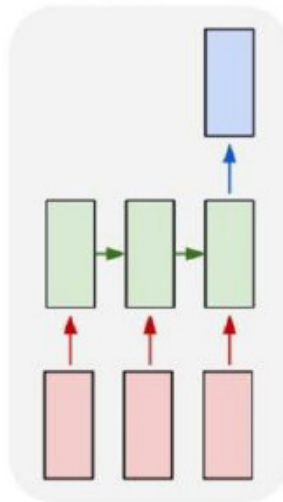
one to one



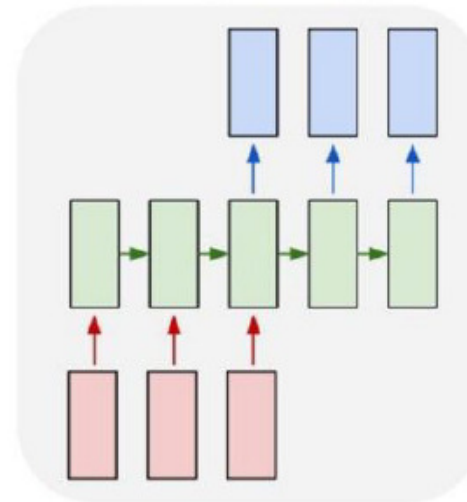
one to many



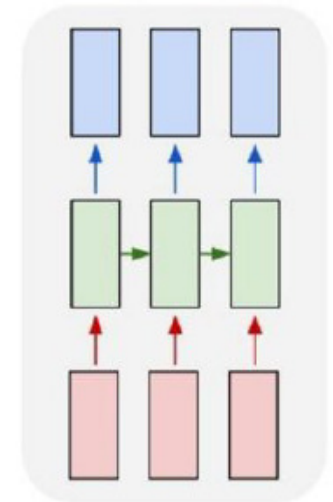
many to one



many to many



many to many

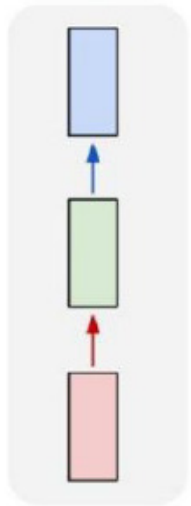


e.g. **Sentiment Classification**  
sequence of words -> sentiment

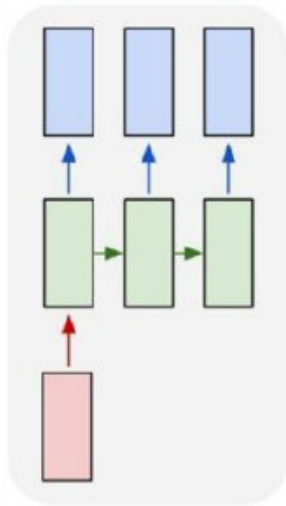
# Recurrent Neural Networks (RNNs)

RNNs can handle

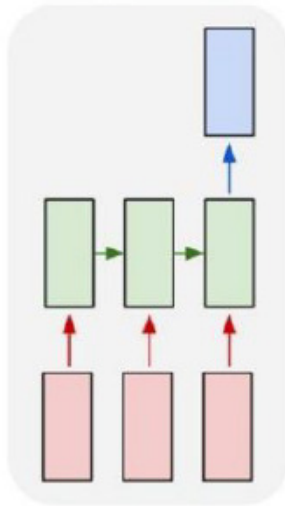
one to one



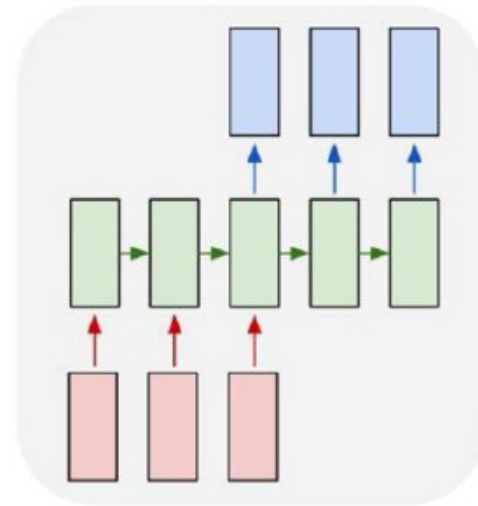
one to many



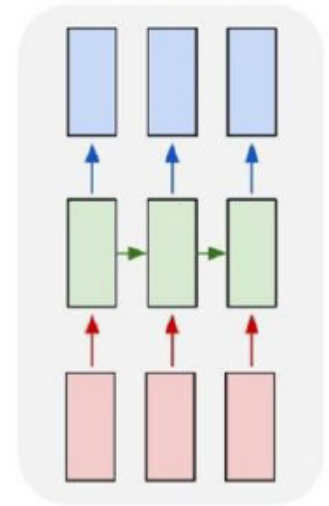
many to one



many to many



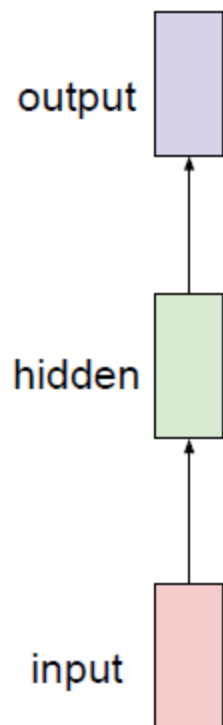
many to many



↖ e.g. **Machine Translation**  
seq of words → seq of words

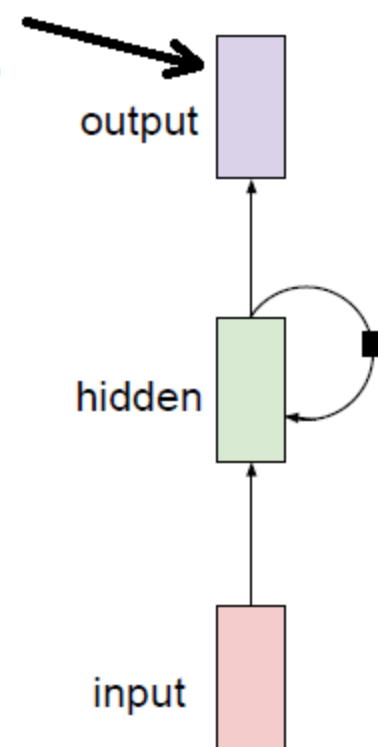
# Recurrent Neural Networks (RNNs)

**Traditional “Feed Forward” Neural Network**



**Recurrent Neural Network**

predict a vector  
at each timestep





# Recurrent Neural Networks

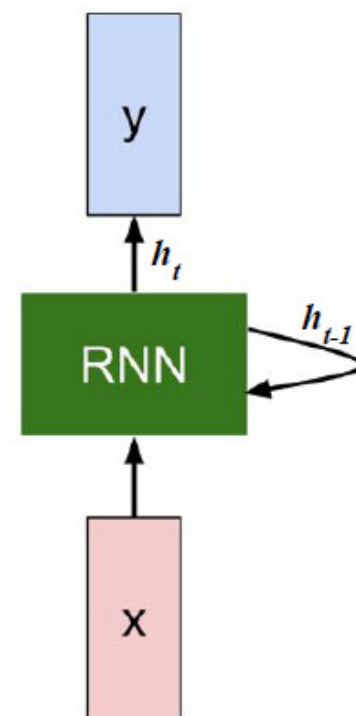
We can process a sequence of vectors  $\mathbf{x}$  by applying a recurrence formula at every time step:

$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state / some function with parameters  $W$

old state

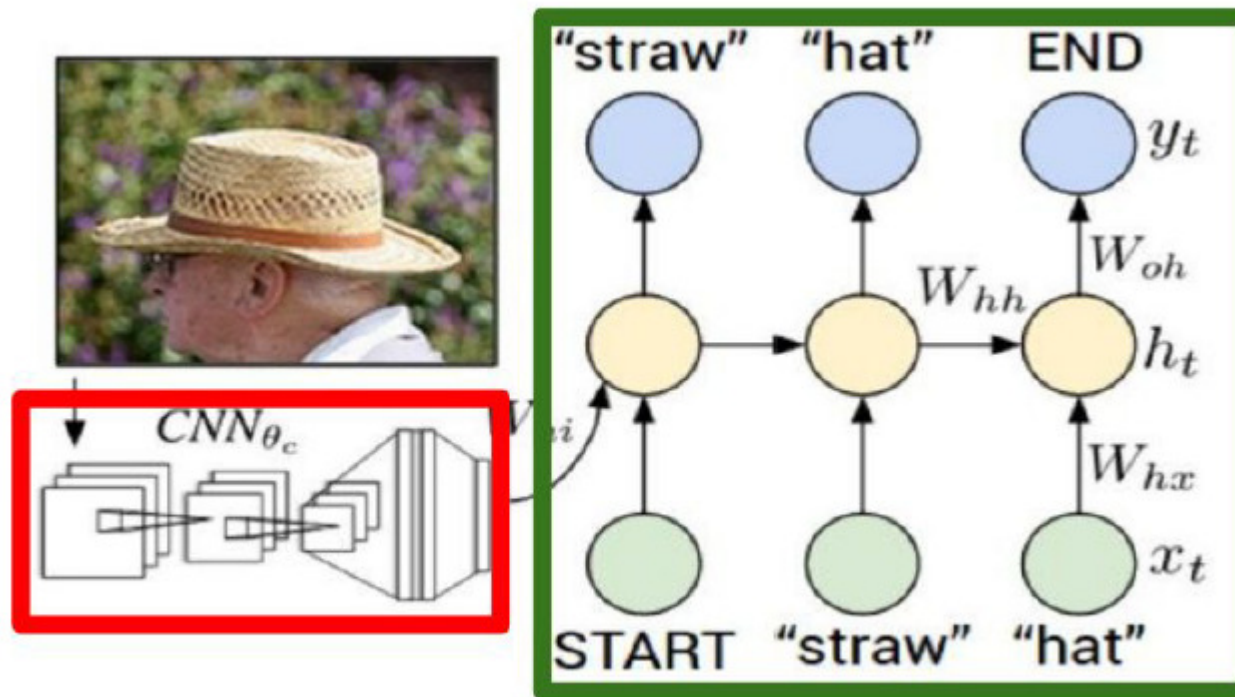
input vector at some time step





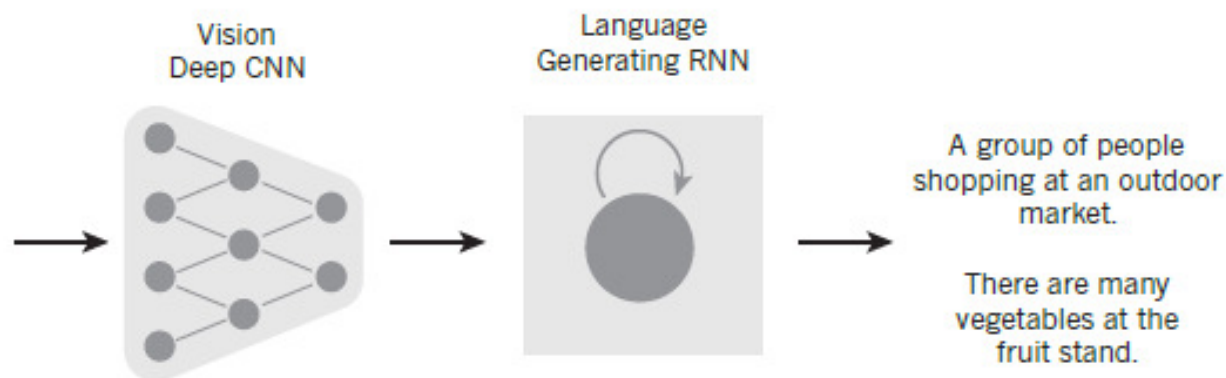
# RNN and CNN Together (*Caption Generation*)

## Recurrent Neural Network



## Convolutional Neural Network

# RNN and CNN Together (*Caption Generation*)



Yann LeCun et al. Nature, Vol. 521, 2015

# RNN and CNN Together (*Caption Generation*)



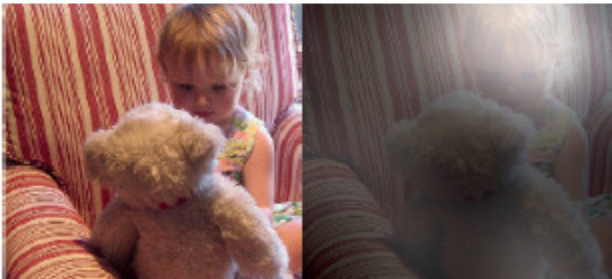
A woman is throwing a **frisbee** in a park.



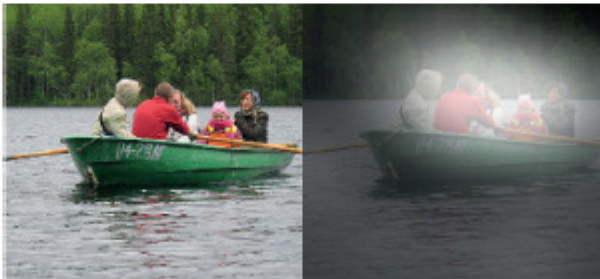
A **dog** is standing on a hardwood floor.



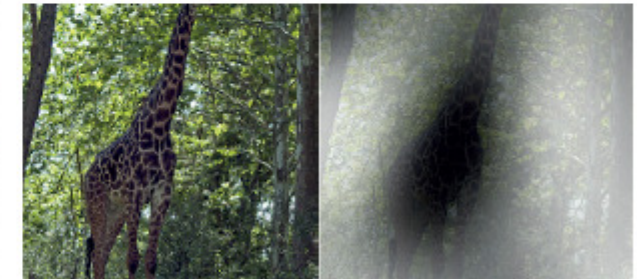
A **stop** sign is on a road with a mountain in the background



A little **girl** sitting on a bed with a teddy bear.



A group of **people** sitting on a boat in the water.



A giraffe standing in a forest with **trees** in the background.

Vinyals, O., Toshev, A., Bengio, S. & Erhan, D. Show and tell: a neural image caption generator. In *Proc. International Conference on Machine Learning* <http://arxiv.org/abs/1502.03044> (2014).

Yann LeCun et al. Nature, Vol. 521, 2015

# CNNs Are a Success!

- The success of CNNs has brought about a revolution in computer vision
- CNNs are now the dominant approach for almost all recognition and detection tasks
- CNNs approach human performance on some tasks!
- Stunning demonstration of CNN + RNN for the generation of image captions (*seen in previous two slides*)



# Some Software

- Many options exist for deep learning
- Main programming languages:
  - Python (<https://www.python.org>)
  - Lua (<https://www.lua.org>)
  - Matlab (<https://www.mathworks.com>)
- Main frameworks:
  - Theano
  - TensorFlow
  - Keras
  - ...more!

# Python Software

- **Anaconda Python** (Python with some pre-loaded libraries) [<https://www.continuum.io/downloads>]
- **Natural Language Toolkit / NLTK** (Python library for natural language processing) [installed with Anaconda]
- **Theano** (Symbolic functions, needed for deep neural networks)
- **Keras** (Layer for arbitrary neural networks, including RNNs/CNNs/etc)
- **Gensim** (Natural Language Processing, word2vec)





## *References:*

~NN and Deep Learning ~ Dr.Qi

~Deep Learning ~ Yann LeCun, Yoshua Bengio & Geoffrey Hinton ~Nature, Vol 521, May 2015

~Imagenet classification with deep convolutional neural networks ~ A.Krizhevsky et al. ~ NIPS 2012

~Very deep convolutional networks for large-scale image recognition ~ K.Simonyan and A Zisserman ~ ICLR 2015

~Going deeper with convolutions ~ C.Szegedy et al. ~ CVPR 2015

~Filters and Convolution ~ <https://en.Wikipedia.org/wiki/Convolution>

~Convolutional Neural Networks & Recurrent Neural Networks: [cs231n.stanford.edu/](http://cs231n.stanford.edu/)