

In-Class Activity 5Source Code:

```
#
#=====
#OOB Errors for Random Forests
#=====
#
#The ``RandomForestClassifier`` is trained using *bootstrap
  aggregation*, where
#each new tree is fit from a bootstrap sample of the training
  observations
#:math:`z_i = (x_i, y_i)`. The *out-of-bag* (OOB) error is the average
  error for
#each :math:`z_i` calculated using predictions from the trees that do
  not
#contain :math:`z_i` in their respective bootstrap sample. This allows
  the
#``RandomForestClassifier`` to be fit and validated whilst being
  trained [1]_.
#
#The example below demonstrates how the OOB error can be measured at
  the
#addition of each new tree during training. The resulting plot allows
  a
#practitioner to approximate a suitable value of ``n_estimators`` at
  which the
#error stabilizes.
#
#.. [1] T. Hastie, R. Tibshirani and J. Friedman, "Elements of
  Statistical
  Learning Ed. 2", p592-593, Springer, 2009.
#=====
#
```

```

import matplotlib.pyplot as plt

from collections import OrderedDict

from sklearn.datasets import make_classification
from sklearn.ensemble import RandomForestClassifier,
    ExtraTreesClassifier

# Reference: scikit-learn.org

print(__doc__)

RANDOM_STATE = 123

# Generate a binary classification dataset.
X, y = make_classification(n_samples=500, n_features=25,
                           n_clusters_per_class=1, n_informative=15,
                           random_state=RANDOM_STATE)

# NOTE: Setting the `warm_start` construction parameter to `True`
        disables

# support for parallelized ensembles but is necessary for tracking the
    OOB

# error trajectory during training.
ensemble_clfs = [
    ("RandomForestClassifier, max_features='sqrt'",
     RandomForestClassifier(warm_start=True, oob_score=True,
                           max_features="sqrt",
                           random_state=RANDOM_STATE)),
    ("RandomForestClassifier, max_features='log2'",
     RandomForestClassifier(warm_start=True, max_features='log2',

```

```

        oob_score=True,
        random_state=RANDOM_STATE)),
    ("RandomForestClassifier, max_features=None",
     RandomForestClassifier(warm_start=True, max_features=None,
                           oob_score=True,
                           random_state=RANDOM_STATE))
]

# Map a classifier name to a list of (<n_estimators>, <error rate>)
# pairs.
error_rate = OrderedDict((label, []) for label, _ in ensemble_clfs)

# Range of `n_estimators` values to explore.
min_estimators = 15
max_estimators = 175

for label, clf in ensemble_clfs:
    for i in range(min_estimators, max_estimators + 1):
        clf.set_params(n_estimators=i)
        clf.fit(X, y)

        # Record the OOB error for each `n_estimators=i` setting.
        oob_error = 1 - clf.oob_score_
        error_rate[label].append((i, oob_error))

# Generate the "OOB error rate" vs. "n_estimators" plot.
for label, clf_err in error_rate.items():
    xs, ys = zip(*clf_err)
    plt.plot(xs, ys, label=label)

plt.xlim(min_estimators, max_estimators)

```

```

plt.title("How an appropriate number of features might be chosen")
plt.xlabel("n_estimators (number of trees)")
plt.ylabel("Out-Of-Bag (OOB) error rate")
plt.legend(loc="upper right")
plt.show()

## Modified from the example written by yhat that can be found here:
#http://blog.yhathq.com/posts/random-forests-in-python.html

from sklearn.ensemble import RandomForestClassifier as RFC
from sklearn.datasets import load_breast_cancer

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

breast = load_breast_cancer()
df = pd.DataFrame(breast.data, columns=breast.feature_names)
df['is_train'] = np.random.uniform(0, 1, len(df)) <= .75
df['species'] = pd.Categorical.from_codes(breast.target,
                                          breast.target_names)

train, test = df[df['is_train']==True], df[df['is_train']==False]
f1 = df.columns[26:27]
f2 = df.columns[25:26]
f3 = df.columns[23:24]
f4 = df.columns[6:7]
f5 = df.columns[10:11]
# features = f1

```

```
# features = f2
# features = f3
# features = f4
# features = f5

forest = RFC(n_jobs=2,n_estimators=50)
y, _ = pd.factorize(train['species'])
forest.fit(train[features], y)

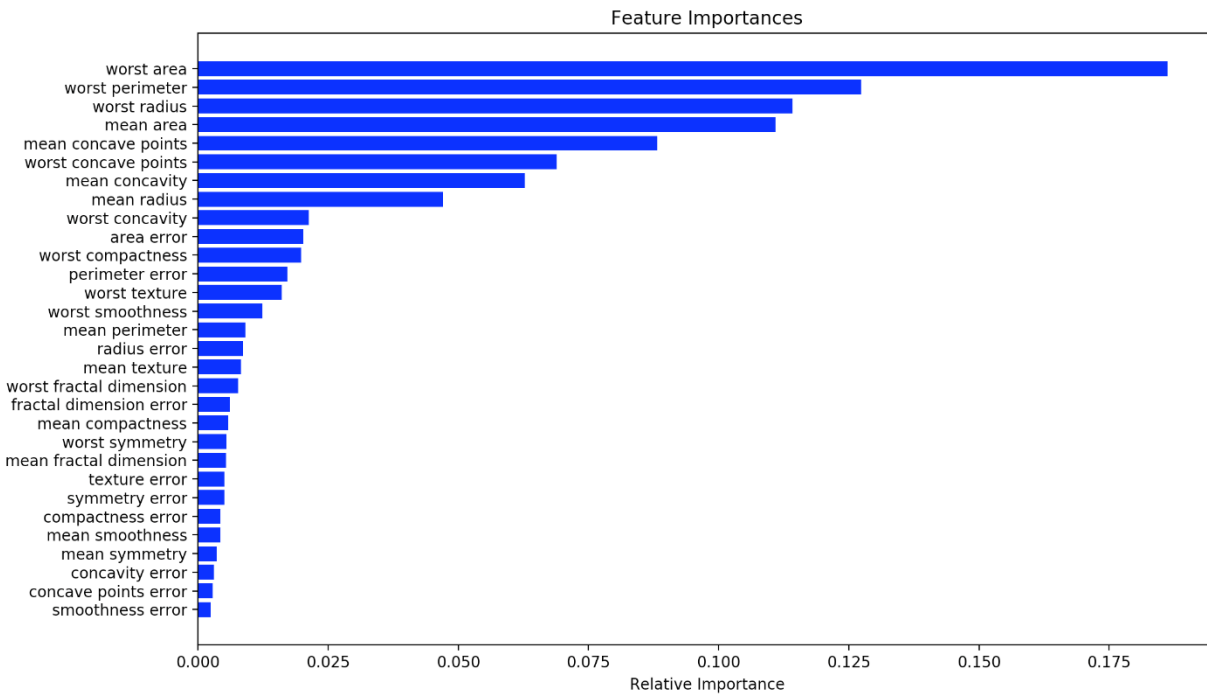
preds = breast.target_names[forest.predict(test[features])]
print (pd.crosstab(index=test['species'], columns=preds,
                    rownames=['actual'], colnames=['preds']))

importances = forest.feature_importances_
indices = np.argsort(importances)

plt.figure(1)
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='b',
          align='center')
plt.yticks(range(len(indices)), features[indices])
plt.xlabel('Relative Importance')
plt.show()
```

Questions:

- 1 We used the breast cancer data set from <http://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Diagnostic%29>.
- 2 The original data set has 32 features.
- 3



- 4 (Removed all the columns except worst area, worst perimeter, worst radius, mean area, and mean concave points. See source code above.)
- 5 The model performed better since we only chose to keep the features that impacted the result (malignant vs. benign) the most, and we didn't remove too many features.